

Data Preprocessing

Machine Learning for Engineering Applications

Fall 2023


Dealing with Missing Data

- **Missing data** from the dataset is very common
 - You will see blanks, NULLs, N/A, etc...
 - There are techniques to deal with missing data
 - Caution: The wrong technique can cause a bad-design model & accuracies
-

- **Pandas!** (Not the cute animal we all love)
- Pandas is a Python library for data analytics
- Pandas is great for importing & managing data in the world of Python
- Pandas is one of the most used libraries for data management in the Machine Learning World!
- Site: <https://pandas.pydata.org/>

pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



Importing and identifying missing data

```
import pandas as pd
from io import StringIO
```

```
csv_data = \
...     '''A,B,C,D
...     1.0,2.0,3.0,4.0
...     5.0,6.0,,8.0
...     10.0,11.0,12.0,'''
```

```
df = pd.read_csv(StringIO(csv_data))
```

```
df
```

	A	B	C	D
0	1.0	2.0	3.0	4.0
1	5.0	6.0	NaN	8.0
2	10.0	11.0	12.0	NaN



Location and count of missing data

```
import pandas as pd
from io import StringIO
```

```
csv_data = \
... '''A,B,C,D
... 1.0,2.0,3.0,4.0
... 5.0,6.0,,8.0
... 10.0,11.0,12.0,'''
```

```
df = pd.read_csv(StringIO(csv_data))
df.isnull().sum()
```

A	0
B	0
C	1
D	1



Numpy Array-like representation of Data Frame


```
import pandas as pd
```

```
csv_data = \  
... '''A,B,C,D  
... 1.0,2.0,3.0,4.0  
... 5.0,6.0,,8.0  
... 10.0,11.0,12.0,'''
```

```
array([[ 1.,  2.,  3.,  4.],  
       [ 5.,  6., nan,  8.],  
       [10., 11., 12., nan]])
```

```
df = pd.read_csv(csv_data)
```

```
df.values
```



Technique of Elimination


Numpy Array-like representation of Data Frame

```
import pandas as pd
```

```
csv_data = \  
... '''A,B,C,D  
... 1.0,2.0,3.0,4.0  
... 5.0,6.0,,8.0  
... 10.0,11.0,12.0,'''
```

	A	B	C	D
0	1.0	2.0	3.0	4.0

```
df = pd.read_csv(csv_data)  
df.dropna(axis=0)
```



- axis=0 means: eliminate rows (samples) with “*nan*” values

Numpy Array-like representation of Data Frame

```
import pandas as pd
```

```
csv_data = \  
... '''A,B,C,D  
... 1.0,2.0,3.0,4.0  
... 5.0,6.0,,8.0  
... 10.0,11.0,12.0,'''
```

```
df = pd.read_csv(csv_data)  
df.dropna(axis=1)
```

	A	B
0	1.0	2.0
1	5.0	6.0
2	10.0	11.0

- axis=1 means: eliminate features (columns) with “*nan*” values

- Threshold operation

`df.dropna(thresh=4)` 

	A	B	C	D
0	1.0	2.0	3.0	4.0

- Removes samples with less than four feature values in the data frame

- Subset operation

`df.dropna(subset=['C'])` 

	A	B	C	D
0	1.0	2.0	3.0	4.0
2	10.0	11.0	12.0	NaN

- Goes into feature 'C' and remove samples with "*nan*"
-

- Danger in removing too many samples
 - Analysis is not reliable or impossible
 - Danger in removing too many features
 - Valuable feature information missing can affect the label classes needed for the machine to learn from
-

Imputing Missing Values

- In statistics, imputation is the process of replacing missing data with substituted values. When substituting for a data point, it is known as "*unit imputation*;" when substituting for a component of a data point, it is known as "*item imputation*."

-- Wikipedia

- Dr. Valles calls it: **Interpolation**
 - There are plenty of interpolation techniques to help in adding/replacing values to your dataset.
-

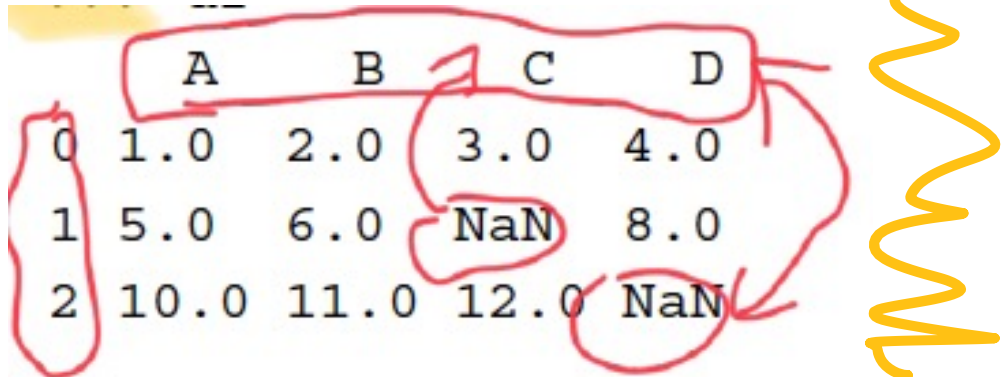
Average/Mean imputation

```
from sklearn.preprocessing import Imputer
```

```
imr = Imputer(missing_values='NaN', strategy='mean', axis=0)  
imr = imr.fit(df.values)  
imputed_data = imr.transform(df.values)  
imputed_data
```

By feature

- Median
- Most frequent
- Less frequent
- Etc



	A	B	C	D
0	1.0	2.0	3.0	4.0
1	5.0	6.0	NaN	8.0
2	10.0	11.0	12.0	NaN

```
array([[ 1.,  2.,  3.,  4.],  
       [ 5.,  6.,  7.5,  8.],  
       [10., 11., 12.,  6.]])
```

Strings & Numerical Values

- Data comes in different **flavors**
 - Some values consists of numerical values: integers, decimals, scientific notation, etc.
 - **Some values consist of characters:** string of characters, single characters, special characters, etc.
 - ML Models evaluate information in a numerical fashion; therefore, we need to “**transform**” string-like values to something that is workable.
-

- **Ordinal features:** values that can be categorized in a sorted or listed in an understandable order.
 - **Example:** Clothing sizes
 - Small < Medium < Large < X-Large < so on
 - These are not numerical
 - However, the order is universally understood
 - Ordinal features can then be arranged in such a way the ML model can “*understand*” the order its learning or the impact of such order.
-

- **Nominal features:** values are understood but do not have an order meaning or impact to its own information.
 - **Example:** Colors
 - Blue < Red < Pink < Grey < so on ???
 - These are not numerical, and no order is understood
 - However, the value of the feature is known to be important
 - Nominal features will have to be transformed to a numeral form that has represents the same information and it must be impactful to the overall learning of the machine.
-

```
import pandas as pd
```

```
df = pd.DataFrame( [  
    ... ['green', 'M', 10.1, 'class1'],  
    ... ['red', 'L', 13.5, 'class2'],  
    ... ['blue', 'XL', 15.3, 'class1']] )  
df.columns = ['color', 'size', 'price', 'classlabel']  
df
```

	color	size	price	classlabel
0	green	M	10.1	class1
1	red	L	13.5	class2
2	blue	XL	15.3	class1

Color: Nominal Feature
Size: Ordinal Feature
Price: Numerical Feature

- **Ordinal features:** values can be assigned!
- **Example:** Clothing Sizes

```
size_mapping = {  
    ... 'XL': 3,  
    ... 'L': 2,  
    ... 'M': 1}  
df['size'] = df['size'].map(size_mapping)  
df
```

	color	size	price	classlabel
0	green	1	10.1	class1
1	red	2	13.5	class2
2	blue	3	15.3	class1

Don't forget the other ordinal feature!


This is transforming
'classlabel' using Numpy

```
import numpy as np
```

```
class_mapping = {label:idx for idx,label in  
    ...enumerate(np.unique(df['classlabel'])) }
```

```
df['classlabel'] = df['classlabel'].map(class_mapping)
```

df



	color	size	price
0	green	1	10.1
1	red	2	13.5
2	blue	3	15.3

classlabel
0
1
0

Don't forget the other ordinal feature!...part 2


```
from sklearn.preprocessing import LabelEncoder
```

```
class_le = LabelEncoder()
```

```
y =  
class_le.fit_transform(df['classlabel'].values)
```

This is transforming
'classlabel' using
Scikit-Learn

y



```
array([0, 1, 0])
```

Encoding

(*Nominal*)

- Encoders in Python are great!... Sort of...
 - They help to quickly take care of the transformation of nominal features
 - The simplest case is the Label Encoder
 - This encoder detects string values
 - Assigns integer values
-

What's the problem with this...?

Example:

```
X = df[['color', 'size', 'price']].values
```

```
color_le = LabelEncoder()
```

```
X[:, 0] = color_le.fit_transform(X[:, 0])
```

X

```
array([[1, 1, 10.1],  
       [2, 2, 13.5],  
       [0, 3, 15.3]], dtype=object)
```

?

blue = 0
green = 1
red = 2

One-Hot Encoding: This will help with the problem below

- This technique will help create temporary (dummy) features to better break the Nominal feature.
 - Binary representations for each feature
 - Expand the number of columns to add all possible differences
-

- **Example**

```
from sklearn.preprocessing import OneHotEncoder
```

```
ohe = OneHotEncoder(categorical_features=[0])  
ohe.fit_transform(X)
```

COLOR

```
array([[ 0. ,  1. ,  0. ,  1. , 10.1],  
       [ 0. ,  0. ,  1. ,  2. , 13.5],  
       [ 1. ,  0. ,  0. ,  3. , 15.3]])
```

Don't Forget...
Partition your
Data

```
from sklearn.model_selection import train_test_split
```

```
X, y = df_wine.iloc[:, 1:].values, df_wine.iloc[:, 0].values
```

```
X_train, X_test, y_train, y_test =\  
    ... train_test_split(X, y,  
    ... test_size=0.3,  
    ... random_state=0,  
    ... stratify=y)
```

Normalization & Standardization

Already covered this one

- Standardized the feature range where the Mean = 0 and Standard Deviation = 1

```
from sklearn.preprocessing import  
StandardScaler
```

```
stdsc = StandardScaler()
```

```
X_train_std = stdsc.fit_transform(X_train)
```

```
X_test_std = stdsc.transform(X_test)
```

$$x_{std}^{(i)} = \frac{x^{(i)} - \mu_x}{\sigma_x}$$

Input	Standardized
0.0	-1.46385
1.0	-0.87831
2.0	-0.29277
3.0	0.29277
4.0	0.87831
5.0	1.46385

Normalization transforms the range of the feature [0,1]

- Knowing the minimum & maximum is crucial
- Therefore, Python has the magical min-max scaling capabilities

```
from sklearn.preprocessing import MinMaxScaler
```


```
mms = MinMaxScaler()
```

```
X_train_norm = mms.fit_transform(X_train)
```

```
X_test_norm = mms.transform(X_test)
```

$$x_{norm}^{(i)} = \frac{x^{(i)} - x_{\min}}{x_{\max} - x_{\min}}$$

- For every data sample in the feature:




- Not all models require features to be normalized or standardize
 - Random Forest, Decision Tree, k -NNs do not really benefit from having its features scaled.
 - Some highly require it due to the design and architecture of the model (*perceptron*)
 - Research will lead you to the right answers!
-

Impactful Features


- **L1** yields to spares feature vectors (many zeros for feature weights)
- **L1** is mainly used as the technique to perform feature selection
- This is useful for *high dimensional datasets*
- The **L1 Penalty**:

$$L1: \| \mathbf{w} \|_1 = \sum_{j=1}^m |w_j|$$



- **L2** is a penalty term that is added to the cost function of the model
- **L2** help to regulate the size of the big weights vs smaller weights
- The penalty strength is the *lambda* term (λ)
- The **L2 Penalty**:

$$L2: \|\mathbf{w}\|_2^2 = \sum_{j=1}^m w_j^2$$



```
from sklearn.linear_model import LogisticRegression
```

```
lr = LogisticRegression(penalty='l1', C=1.0)
```

```
lr.fit(X_train_std, y_train)
```

```
print('Training accuracy:', lr.score(X_train_std, y_train))
```

```
Training accuracy: 1.0
```

```
print('Test accuracy:', lr.score(X_test_std, y_test))
```

```
Test accuracy: 1.0
```

```
lr.coef_
```

```
array([[1.24559337, 0.18041967, 0.74328894, -1.16046277, 0., 0., 1.1678711,  
0., 0., 0., 0., 0.54941931, 2.51017406],
```

```
[-1.53720749, -0.38727002, -0.99539203, 0.3651479, -0.0596352 , 0.,  
0.66833149, 0., 0., -1.9346134, 1.23297955, 0., -2.23135027],
```

```
[ 0.13579227, 0.16837686, 0.35723831, 0., 0., 0., -2.43809275, 0., 0.,  
1.56391408, -0.81933286, -0.49187817, 0.]])
```

Feature Selection & Extraction

- One of the ways to avoid overfitting: **Dimensionality reduction**
 - **Two types:**
 - Feature Selection
 - Feature Extraction
 - **Feature Selection:** generate a subset of features
 - **Feature Extraction:** create a new feature from features
-

Next Assignment
