



PCA, LDA, Kernel PCA, and t-SNE/UMAPs

Machine Learning for Engineering Applications

Fall 2023


Data Compression

- **Principal Component Analysis (PCA)**
 - **Unsupervised** data compression
 - **Linear Discriminant Analysis (LDA)**
 - **Supervised** dimensionality reduction
 - **Kernel Principal Component Analysis (KPCA)**
 - **Non-linear** data compression approach
-

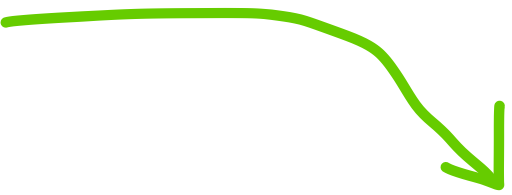
PCA

- In the context of dimensionality reduction:
 - feature extraction is the approach to data compression
 - **Goal**: maintaining most of the relevant information
 - **PCAs** are widely used when you have a common problem: **too many features!!!**
 - This will force you to perform a dimensional reduction without losing the information of all of the dataset
-


- A transformation helps to “*encapsulate*” the meaning of the data
- The PC components help to keep the integrity of the reduction of features
- The PCs need to be orthogonal to obtain a proper transformation of the dataset as its being reduced




- Original feature space


$$\mathbf{x} = [x_1, x_2, \dots, x_d], \quad \mathbf{x} \in \mathbb{R}^d$$

- Transformation operation that consider the reduction of a *d-dimensional* space to a *k-dimensional* space


$$\downarrow \mathbf{x}W, \quad W \in \mathbb{R}^{d \times k}$$

- The compressed feature space


$$\mathbf{z} = [z_1, z_2, \dots, z_k], \quad \mathbf{z} \in \mathbb{R}^k$$

7-STEPS to PCA



1. Standardize the d -dimensional dataset.
 2. Construct the covariance matrix.
 3. Decompose the covariance matrix into its eigenvectors and eigenvalues.
 4. Sort the eigenvalues by decreasing order to rank the corresponding eigenvectors.
-

5. Select k eigenvectors which correspond to the k largest eigenvalues, where k is the dimensionality of the new feature subspace ($k \leq d$).
 6. Construct a projection matrix **W** from the "top" k -eigenvectors.
 7. Transform the d -dimensional input dataset **X** using the projection matrix W to obtain the new k -dimensional feature subspace.
-

- Wines!!!
- The dataset has 13 features....who knew...
- Do you need all the features?
- Which features are more important?
- What if you do not know anything about wines?



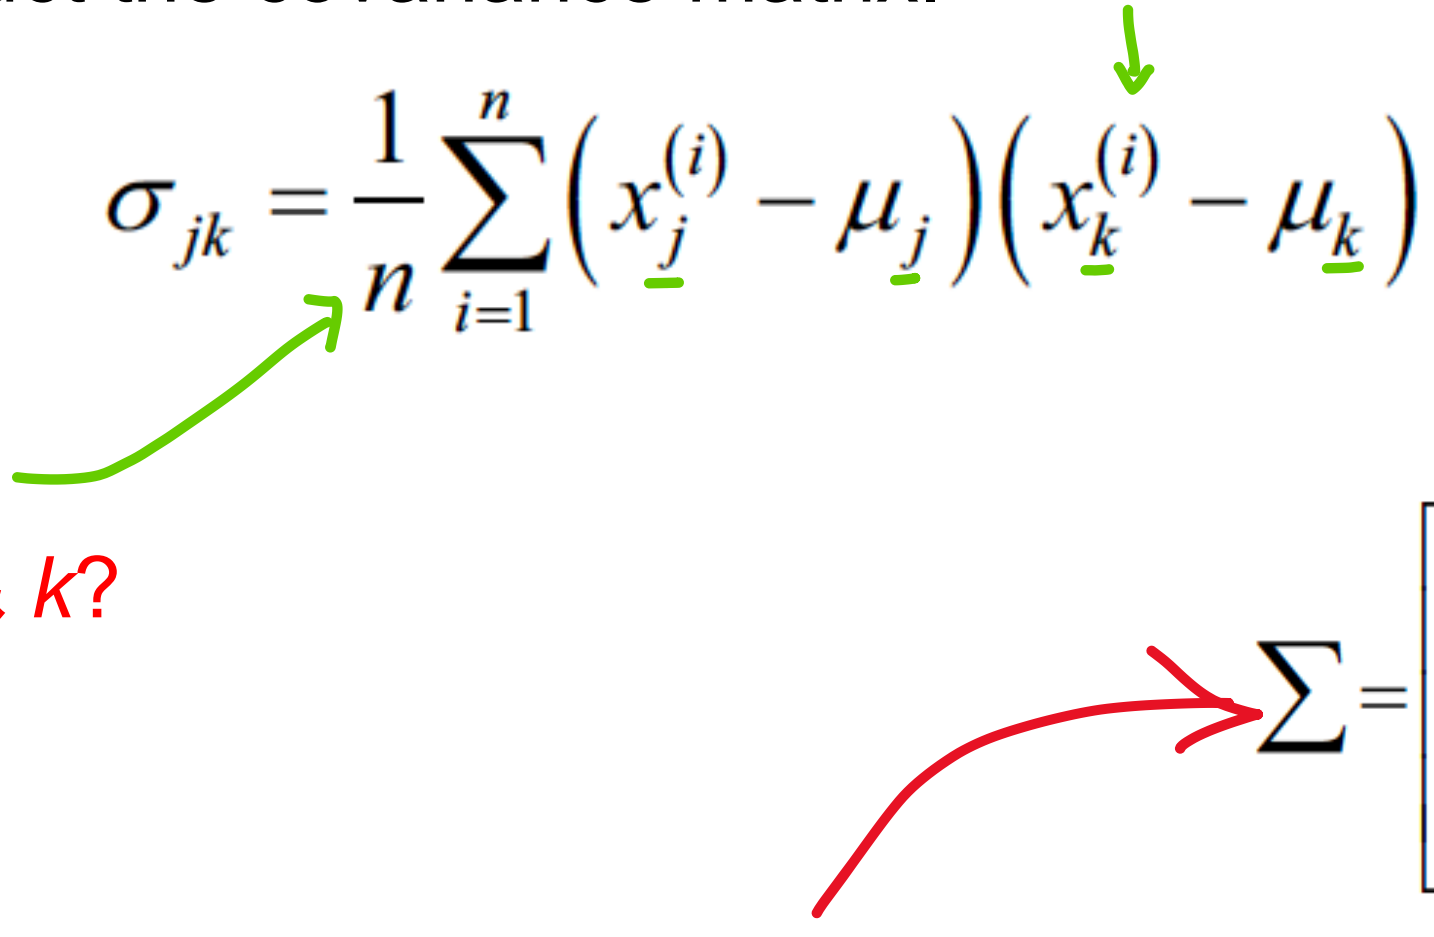
1. Standardize the d -dimensional dataset.

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```
df_wine =
pd.read_csv('https://archive.ics.uci.edu/ml/'machine-
learning-databases/wine/wine.data', header=None)
X, y = df_wine.iloc[:, 1:].values, df_wine.iloc[:, 0].values
X_train, X_test, y_train, y_test =
... train_test_split(X, y, test_size=0.3, stratify=y,
... random_state=0)
```

```
sc = StandardScaler()
X_train_std = sc.fit_transform(X_train)
X_test_std = sc.transform(X_test)
```

2. Construct the covariance matrix.

$$\sigma_{jk} = \frac{1}{n} \sum_{i=1}^n \left(x_{\underline{j}}^{(i)} - \mu_{\underline{j}} \right) \left(x_{\underline{k}}^{(i)} - \mu_{\underline{k}} \right)$$


What's n ?

What's j & k ?

What's i ?

$$\Sigma = \begin{bmatrix} \sigma_1^2 & \sigma_{12} & \sigma_{13} \\ \sigma_{21} & \sigma_2^2 & \sigma_{23} \\ \sigma_{31} & \sigma_{32} & \sigma_3^2 \end{bmatrix}$$

Covariance matrix for a 3-feature dataset

3. Decompose the covariance matrix into its eigenvectors and eigenvalues.

- The eigenvectors become the Principal Components of your dataset
- The eigenvalues are the level of importance to their corresponding eigenvectors
- **Review:** Eigenvectors must satisfy, $\Sigma \mathbf{v} = \lambda \mathbf{v}$


3. Decompose the covariance matrix into its eigenvectors and eigenvalues

```
sc = StandardScaler()  
X_train_std = sc.fit_transform(X_train)  
X_test_std = sc.transform(X_test)
```

```
import numpy as np  
cov_mat = np.cov(X_train_std.T)  
eigen_vals, eigen_vecs = np.linalg.eig(cov_mat)  
print('\nEigenvalues \n%s' % eigen_vals)
```

Eigenvalues

```
[ 4.84274532  2.41602459  1.54845825  0.96120438  0.84166161  
 0.6620634    0.51828472  0.34650377  0.3131368   0.10754642  
 0.21357215   0.15362835  0.1808613 ]
```



4. Sort the eigenvalues by decreasing order to rank the corresponding eigenvectors.

- This is a highly step to not just perform in the code...take time to analyze it as well!
- **The importance:** you will visually see the impact of the features in your dataset
- What's going on for this step? Checking eigenvalues to the aggregated sum...

$$\frac{\lambda_j}{\sum_{j=1}^d \lambda_j}$$

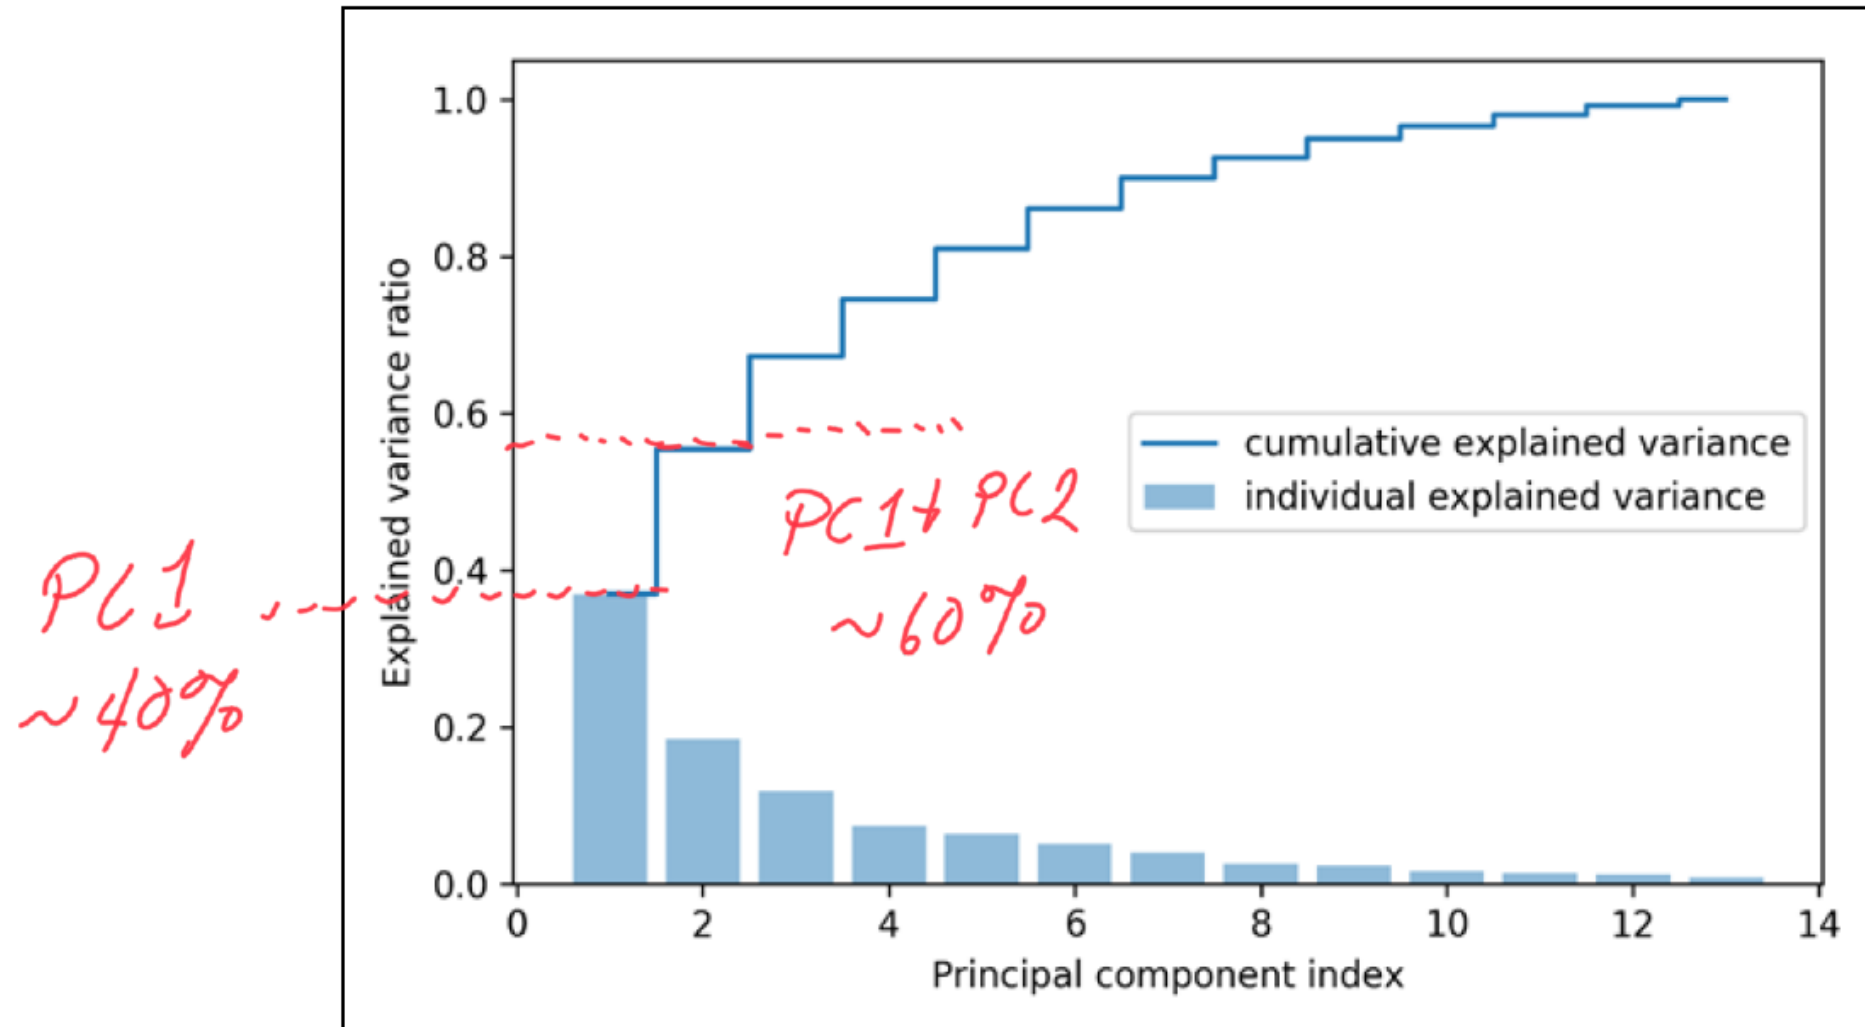
4. Sort the eigenvalues by decreasing order to rank the corresponding eigenvectors.

```
tot = sum(eigen_vals)
var_exp = [(i / tot) for i in sorted(eigen_vals,
                                     ... reverse=True)]

cum_var_exp = np.cumsum(var_exp)

# plot code of the next slide
import matplotlib.pyplot as plt
plt.bar(range(1,14), var_exp, alpha=0.5, align='center',
          ... label='individual explained variance')
plt.step(range(1,14), cum_var_exp, where='mid',
          ... label='cumulative explained variance')
plt.ylabel('Explained variance ratio')
plt.xlabel('Principal component index')
plt.legend(loc='best')
plt.show()
```

4. Sort the eigenvalues by decreasing order to rank the corresponding eigenvectors.



5. Select k eigenvectors which correspond to the k largest eigenvalues, where k is the dimensionality of the new feature subspace ($k \leq d$).

- Create pairs to sort the eigenvector by the eigenvalues
 - **First Element:** Eigenvalue
 - **2nd Element:** The associated eigenvector
 - (<<eigenvalue object>>, <<eigenvector object>>)
-

5. Select k eigenvectors which correspond to the k largest eigenvalues, where k is the dimensionality of the new feature subspace ($k \leq d$).

```
eigen_pairs = [(np.abs(eigen_vals[i]), eigen_vecs[:, i])  
for i in range(len(eigen_vals))]  
  
# Sort the (eigenvalue, eigenvector) tuples from high to low  
eigen_pairs.sort(key=lambda k: k[0], reverse=True)
```

6. Construct a projection matrix **W** from the "top" k -eigenvectors.

- Most practitioners pick the top 2 eigenvectors
 - Mainly: See if you can get away the simplest amount of PC units
 - There's a "*high*" chance that your top 2 can cover > 50% of features
-

6. Construct a projection matrix \mathbf{W} from the "top" k -eigenvectors.

```
w = np.hstack((eigen_pairs[0][1][:, np.newaxis],  
               ... eigen_pairs[1][1][:, np.newaxis]))  
print('Matrix W:\n', w)
```

Matrix W:

```
[[-0.13724218  0.50303478]  
 [ 0.24724326  0.16487119]  
 [-0.02545159  0.24456476]  
 [ 0.20694508 -0.11352904]  
 [-0.15436582  0.28974518]  
 [-0.39376952  0.05080104]  
 [-0.41735106 -0.02287338]  
 [ 0.30572896  0.09048885]  
 [-0.30668347  0.00835233]  
 [ 0.07554066  0.54977581]  
 [-0.32613263 -0.20716433]  
 [-0.36861022 -0.24902536]  
 [-0.29669651  0.38022942]]
```

Handwritten annotations:

- pair idx (with arrows pointing to the first two rows of the matrix)
- index of the pair (with an arrow pointing to the first two rows of the matrix)
- all rows for new column (with an arrow pointing to the first two rows of the matrix)

7. Transform the d -dimensional input dataset \mathbf{X} using the projection matrix \mathbf{W} to obtain the new k -dimensional feature subspace.

- Transformation: $\mathbf{X}' = \mathbf{X}\mathbf{W}$

```
X_train_pca = X_train_std.dot(w)
```

```
colors = ['r', 'b', 'g']
```

```
markers = ['s', 'x', 'o']
```

```
for l, c, m in zip(np.unique(y_train), colors, markers):
```

```
    ... plt.scatter(X_train_pca[y_train==l, 0],
```

```
    ... X_train_pca[y_train==l, 1],
```


```
    ... c=c, label=l, marker=m)
```

```
plt.xlabel('PC 1')
```

```
plt.ylabel('PC 2')
```

```
plt.legend(loc='lower left')
```


```
plt.show()
```





The scikit-learn implementation using Logistic Regression:

```
from sklearn.linear_model import LogisticRegression
from sklearn.decomposition import PCA
pca = PCA(n_components=2)
lr = LogisticRegression()
X_train_pca = pca.fit_transform(X_train_std)
X_test_pca = pca.transform(X_test_std)
lr.fit(X_train_pca, y_train) #assuming you did get y_train @ this point as well
plot_decision_regions(X_train_pca, y_train, classifier=lr)
plt.xlabel('PC 1')
plt.ylabel('PC 2')
plt.legend(loc='lower left')
plt.show()

# did not show the fancy coloring code for the output
# see textbook for that part.
```






```
plot_decision_regions(X_test_pca, y_test, classifier=lr)
```

LDA

- In the context of feature extraction:
 - Helps not to overfit the model
 - **This is a supervised method**
 - **LDAs** are widely used when you have a common problem: **too many features!!!**
 - This will force you to perform a dimensional reduction without losing the information of all the dataset using its class labels and
-

- **The big assumptions:**
- The data is normally distributed
- Features are independent from each other
- Covariance matrices for all classes are identical



7-STEPS to LDA



1. Standardize the d -dimensional dataset (d is the number of features).
 2. For each class, compute the d -dimensional mean vector.
 3. Construct the between-class scatter matrix \mathbf{S}_B and the within-class scatter matrix \mathbf{S}_W .
 4. Compute the eigenvectors and corresponding eigenvalues of the matrix $\mathbf{S}_W^{-1}\mathbf{S}_B$
-

5. Sort the eigenvalues by decreasing order to rank the corresponding eigenvectors.
 6. Choose the k eigenvectors that correspond to the k largest eigenvalues to construct a $d \times k$ -dimensional transformation matrix **W**; the eigenvectors are the columns of this matrix.
 7. Project the samples onto the new feature subspace using the transformation matrix **W**.
-

1. Standardize the d -dimensional dataset (d is the number of features).



2. For each class, compute the d -dimensional mean vector.

$$\mathbf{m}_i = \frac{1}{n_i} \sum_{\mathbf{x} \in D_i} \mathbf{x}_m \quad \longrightarrow \quad \mathbf{m}_i = \begin{bmatrix} \mu_{i, \text{alcohol}} \\ \mu_{i, \text{malic acid}} \\ \vdots \\ \mu_{i, \text{proline}} \end{bmatrix} \quad i \in \{1, 2, 3\}$$

A blue arrow points from the summation term \mathbf{x}_m in the first equation to the vector \mathbf{m}_i in the second equation.

Handwritten in red:

$$\mathbf{m}_i = \begin{bmatrix} \mu_{i, ft.1} \\ \mu_{i, ft.2} \\ \vdots \\ \mu_{i, ft.\alpha} \end{bmatrix} \quad i \in \{1 : \# \text{ of classes}\}$$

2. For each class, compute the d -dimensional mean vector.

```
np.set_printoptions(precision=4)
mean_vecs = []
for label in range(1, 4):
    mean_vecs.append(np.mean(
        ... X_train_std[y_train==label], axis=0))
    ... print('MV %s: %s\n' % (label, mean_vecs[label-1]))
print('MV %s: %s\n' % (label, mean_vecs[label-1]))
```

```
MV 1: [ 0.9066 -0.3497 0.3201 -0.7189 0.5056 0.8807 0.9589 -0.5516
       0.5416 0.2338 0.5897 0.6563 1.2075]
```

```
MV 2: [-0.8749 -0.2848 -0.3735 0.3157 -0.3848 -0.0433 0.0635 - 0.0946
       0.0703 -0.8286 0.3144 0.3608 -0.7253]
```

```
MV 3: [ 0.1992 0.866 0.1682 0.4148 -0.0451 -1.0286 -1.2876 0.8287
       -0.7795 0.9649 -1.209 -1.3622 -0.4013]
```

3. Construct the **between-class scatter** matrix \mathbf{S}_B and the **within-class scatter** matrix \mathbf{S}_W .

- The within-class scatter matrix:

$$\mathbf{S}_W = \sum_{i=1}^c \mathbf{S}_i \quad \text{for all classes}$$

$$\mathbf{S}_i = \sum_{\mathbf{x} \in D_i} (\mathbf{x} - \mathbf{m}_i)(\mathbf{x} - \mathbf{m}_i)^T \quad \text{for class } i$$

3. Construct the **between-class scatter** matrix \mathbf{S}_B and the **within-class scatter** matrix \mathbf{S}_W .

- The within-class scatter matrix:

```
d = 13 # number of features
S_W = np.zeros((d, d))
for label, mv in zip(range(1, 4), mean_vecs):
    ... class_scatter = np.zeros((d, d))
for row in X_train_std[y_train == label]:
    ... row, mv = row.reshape(d, 1), mv.reshape(d, 1)
    ... class_scatter += (row - mv).dot((row - mv).T)
    ... S_W += class_scatter
print('Class label distribution: %s' % np.bincount(y_train)[1:])
```

Class label distribution: [41 50 33]... distribution is unbalanced

3. Construct the **between-class scatter** matrix \mathbf{S}_B and the **within-class scatter** matrix \mathbf{S}_W .

- The between-class scatter matrix:

$$\mathbf{S}_B = \sum_{i=1}^c n_i (\mathbf{m}_i - \mathbf{m})(\mathbf{m}_i - \mathbf{m})^T$$

- \mathbf{m} : the overall average that includes samples from other classes

3. Construct the **between-class scatter** matrix \mathbf{S}_B and the **within-class scatter** matrix \mathbf{S}_W .

- The between-class scatter matrix:

```
mean_overall = np.mean(X_train_std, axis=0)
d = 13 # number of features
S_B = np.zeros((d, d))
for i, mean_vec in enumerate(mean_vecs):
    ... n = X_train[y_train == i + 1, :].shape[0]
    ... mean_vec = mean_vec.reshape(d, 1) # make column vector
    ... mean_overall = mean_overall.reshape(d, 1)
    ... S_B += n * (mean_vec - mean_overall).dot(
    ... (mean_vec - mean_overall).T)
```

4. Compute the eigenvectors and corresponding eigenvalues of the matrix $S_W^{-1}S_B$

```
eigen_vals, eigen_vecs =\  
... np.linalg.eig(np.linalg.inv(S_W).dot(S_B))
```

5. Sort the eigenvalues by decreasing order to rank the corresponding eigenvectors.

```
eigen_pairs = [(np.abs(eigen_vals[i]), eigen_vecs[:,i])
                ... for i in range(len(eigen_vals))]
```

```
eigen_pairs = sorted(eigen_pairs,  
    ... key=lambda k: k[0], reverse=True)
```



```
print('Eigenvalues in descending order:\n')  
for eigen_val in eigen_pairs:  
    print(eigen_val[0])
```

Eigenvalues in descending order:

349.617808906

172.76152219

3.78531345125e-14

2.11739844822e-14

1.51646188942e-14

1.51646188942e-14

1.35795671405e-14

1.35795671405e-14

7.58776037165e-15

5.90603998447e-15

5.90603998447e-15


2.25644197857e-15

0.0

Print out the LD magnitudes

```
tot = sum(eigen_vals.real)
discr = [(i / tot) for i in sorted(eigen_vals.real,
reverse=True)]
cum_discr = np.cumsum(discr)

plt.bar(range(1, 14), discr, alpha=0.5, align='center',
... label='individual "discriminability"')
plt.step(range(1, 14), cum_discr, where='mid',
... label='cumulative "discriminability"')
plt.ylabel('"discriminability" ratio')
plt.xlabel('Linear Discriminants')
plt.ylim([-0.1, 1.1])
plt.legend(loc='best')
plt.show()
```



6. Choose the k eigenvectors that correspond to the k largest eigenvalues to construct a $d \times k$ -dimensional transformation matrix **W**; the eigenvectors are the columns of this matrix.

- The first two are the main LDs

```
w = np.hstack((eigen_pairs[0][1][:, np.newaxis].real,  
               ... eigen_pairs[1][1][:, np.newaxis].real))
```

```
print('Matrix W:\n', w)
```


Matrix W:

```
[[ -0.1481  -0.4092]  
[  0.0908  -0.1577]  
[ -0.0168  -0.3537]  
[  0.1484   0.3223]  
[ -0.0163  -0.0817]  
[  0.1913   0.0842]  
[ -0.7338   0.2823]  
[ -0.075   -0.0102]  
[  0.0018   0.0907]  
[  0.294   -0.2152]  
[ -0.0328   0.2747]  
[ -0.3547  -0.0124]  
[ -0.3915  -0.5958]]
```

7. Project the samples onto the new feature subspace using the transformation matrix **W**.


- The projection to new subspace: $X' = XW$

```
X_train_lda = X_train_std.dot(w)
colors = ['r', 'b', 'g']
markers = ['s', 'x', 'o']
for l, c, m in zip(np.unique(y_train), colors, markers):
    ... plt.scatter(X_train_lda[y_train==l, 0],
    ... X_train_lda[y_train==l, 1] * (-1),
    ... c=c, label=l, marker=m)
plt.xlabel('LD 1')
plt.ylabel('LD 2')
plt.legend(loc='lower right')
plt.show()
```



The scikit-version:



```
lr = LogisticRegression()  
lr = lr.fit(X_train_lda, y_train)  
plot_decision_regions(X_train_lda, y_train,  
                        ... classifier=lr)  
  
plt.xlabel('LD 1')  
plt.ylabel('LD 2')  
plt.legend(loc='lower left')  
plt.show()
```

The test split data:

```
X_test_lda = lda.transform(X_test_std)
plot_decision_regions(X_test_lda, y_test,
                      ... classifier=lr)
```

```
plt.xlabel('LD 1')
plt.ylabel('LD 2')
plt.legend(loc='lower left')
plt.show()
```



Kernel PCA

- Kernel to the rescue for **non-linear problems** (*higher dimension*)
- The transformation for KPCA is the key:

$$\phi : \mathbb{R}^d \rightarrow \mathbb{R}^k \quad (k \gg d)$$

$$\mathbf{x} = [x_1, x_2]^T \longrightarrow x \in \mathbb{R}^2$$

$$\downarrow \phi$$

$$\mathbf{z} = [x_1^2, \sqrt{2x_1x_2}, x_2^2]^T \longrightarrow$$

This can be different!

- **Step 1:** Dimensional Transformation
- **Step 2:** Do all the PCA steps!

- Covariance (*review*):

$$\sigma_{jk} = \frac{1}{n} \sum_{i=1}^n \left(x_j^{(i)} - \mu_j \right) \left(x_k^{(i)} - \mu_k \right)$$

- Covariance Matrix (*review*):

$$\Sigma = \frac{1}{n} \sum_{i=1}^n \mathbf{x}^{(i)} \mathbf{x}^{(i)T}$$

- **Step 3:** KPCA Covariance Matrix:

$$\Sigma = \frac{1}{n} \sum_{i=1}^n \phi(\mathbf{x}^{(i)}) \phi(\mathbf{x}^{(i)})^T$$

- Eigenvectors:

$$\Sigma \mathbf{v} = \lambda \mathbf{v}$$

- Eigenvectors:

$$\Sigma \mathbf{v} = \lambda \mathbf{v}$$

$$\Rightarrow \frac{1}{n} \sum_{i=1}^n \phi(\mathbf{x}^{(i)}) \phi(\mathbf{x}^{(i)})^T \mathbf{v} = \lambda \mathbf{v}$$

$$\Rightarrow \mathbf{v} = \frac{1}{n\lambda} \sum_{i=1}^n \phi(\mathbf{x}^{(i)}) \phi(\mathbf{x}^{(i)})^T \mathbf{v}$$

$$= \frac{1}{n} \sum_{i=1}^n \mathbf{a}^{(i)} \phi(\mathbf{x}^{(i)})$$

$\mathbf{a}^{(i)} = \frac{\phi(\mathbf{x}^{(i)})^T}{\lambda}$

- **Kernel Function:**

$$\kappa\left(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}\right)=\phi\left(\mathbf{x}^{(i)}\right)^T \phi\left(\mathbf{x}^{(j)}\right)$$

- This is more useful than trying to figure out the Eigen pairs
 - The kernel PCA are samples are already projected (**transformed**) to the PC components
 - There are different flavors of kernel functions
-

- **Polynomial Kernel:**


$$K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \left(\mathbf{x}^{(i)T} \mathbf{x}^{(j)} + \theta \right)^p$$

- **Hyperbolic Tangent (sigmoid) kernel:**

$$K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \tanh\left(\eta \mathbf{x}^{(i)T} \mathbf{x}^{(j)} + \theta\right)$$

- **Radial Basis Function (Gaussian) kernel:**

$$K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \exp\left(-\frac{\|\mathbf{x}^{(i)} - \mathbf{x}^{(j)}\|^2}{2\sigma^2}\right)$$


$$\gamma = \frac{1}{2\sigma}$$

- Half Moon Example:
- Authors: 2 Scenarios (homebrew vs. PCA)


```
from sklearn.datasets import make_moons

X, y = make_moons(n_samples=100,
                  ... random_state=123)

plt.scatter(X[y==0, 0], X[y==0, 1],
            ... color='red', marker='^', alpha=0.5)

plt.scatter(X[y==1, 0], X[y==1, 1],
            ... color='blue', marker='o', alpha=0.5)



plt.show()
```



- Normal PCA

```
from sklearn.decomposition import PCA
scikit_pca = PCA(n_components=2)
X_spca = scikit_pca.fit_transform(X)
fig, ax = plt.subplots(nrows=1,ncols=2, figsize=(7,3))
ax[0].scatter(X_spca[y==0, 0], X_spca[y==0, 1],
              ... color='red', marker='^', alpha=0.5)
ax[0].scatter(X_spca[y==1, 0], X_spca[y==1, 1],
              ... color='blue', marker='o', alpha=0.5)
ax[1].scatter(X_spca[y==0, 0], np.zeros((50,1))+0.02,
              ... color='red', marker='^', alpha=0.5)
ax[1].scatter(X_spca[y==1, 0], np.zeros((50,1))-0.02,
              ... color='blue', marker='o', alpha=0.5)
ax[0].set_xlabel('PC1')
ax[0].set_ylabel('PC2')
ax[1].set_ylim([-1, 1])
ax[1].set_yticks([])
ax[1].set_xlabel('PC1')
plt.show()
```

- **Normal** PCA Results





- Homebrew RBF Kernel PCA


```
from sklearn.decomposition import PCA
X_kpca = rbf_kernel_pca(X, gamma=15, n_components=2)

fig, ax = plt.subplots(nrows=1,ncols=2, figsize=(7,3))
ax[0].scatter(X_kpca[y==0, 0], X_kpca[y==0, 1],
              ... color='red', marker='^', alpha=0.5)
ax[0].scatter(X_kpca[y==1, 0], X_kpca[y==1, 1],
              ... color='blue', marker='o', alpha=0.5)
ax[1].scatter(X_kpca[y==0, 0], np.zeros((50,1))+0.02,
              ... color='red', marker='^', alpha=0.5)
ax[1].scatter(X_kpca[y==1, 0], np.zeros((50,1))-0.02,
              ... color='blue', marker='o', alpha=0.5)
ax[0].set_xlabel('PC1')
ax[0].set_ylabel('PC2')
ax[1].set_ylim([-1, 1])
ax[1].set_yticks([])
ax[1].set_xlabel('PC1')
plt.show()
```



- Homebrew RBF Kernel PCA Results





- Concentric Circles Example:
- Authors: 2 Scenarios (homebrew vs. PCA)



- Normal PCA Results



- Homebrew RBF Kernel PCA Results



Scikit-Learn Implementation

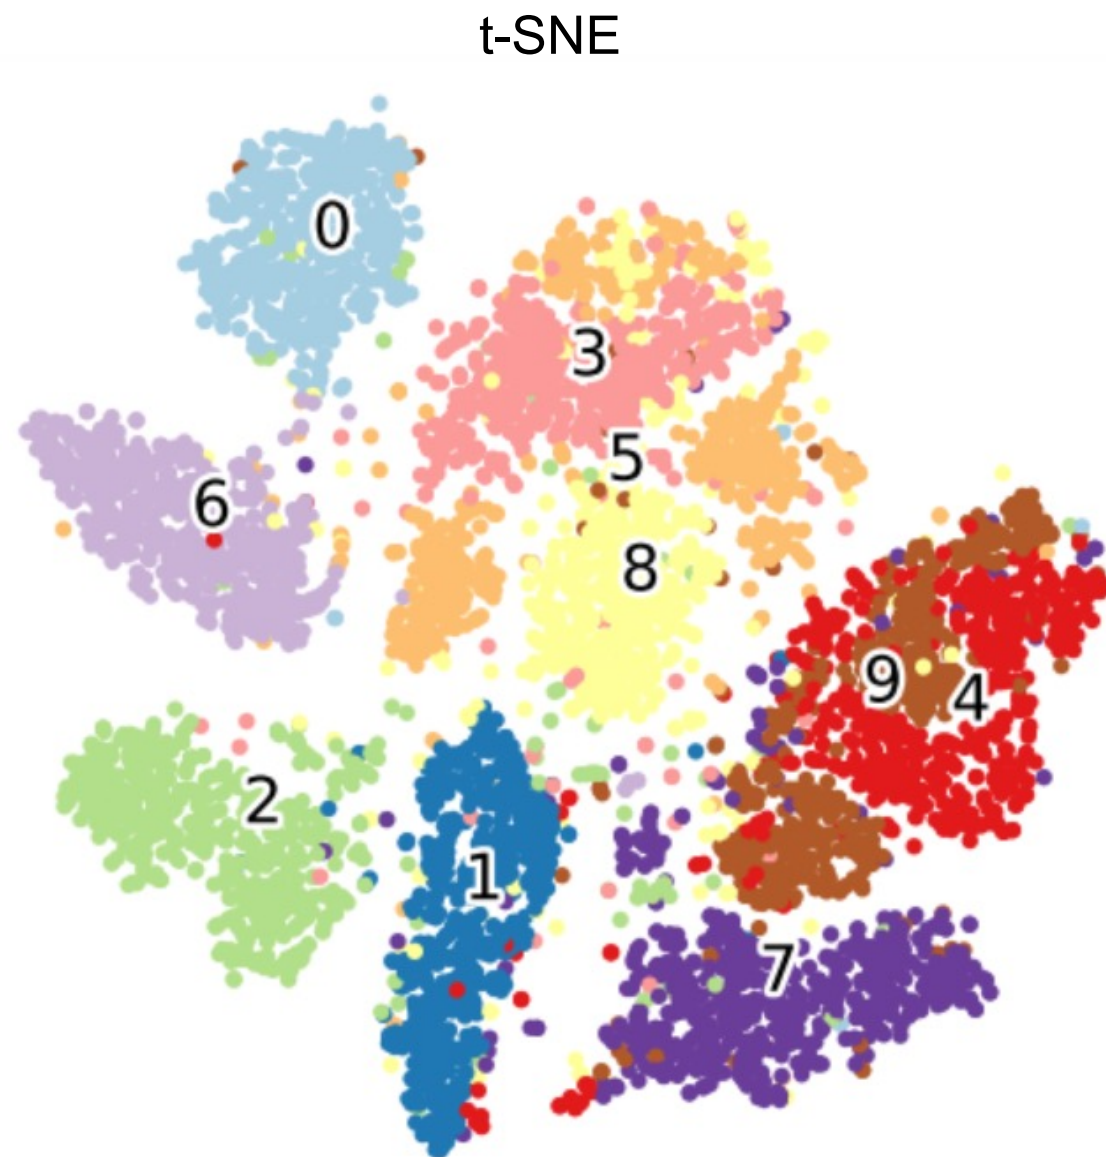

```
from sklearn.decomposition import KernelPCA
X, y = make_moons(n_samples=100, random_state=123)
scikit_kpca = KernelPCA(n_components=2,
                        ... kernel='rbf', gamma=15)
X_skernpca = scikit_kpca.fit_transform(X)

plt.scatter(X_skernpca[y==0, 0], X_skernpca[y==0, 1],
            ... color='red', marker='^', alpha=0.5)
plt.scatter(X_skernpca[y==1, 0], X_skernpca[y==1, 1],
            ... color='blue', marker='o',

alpha=0.5)
plt.xlabel('PC1')
plt.ylabel('PC2')
plt.show()
```

T-distributed Stochastic Neighbor Embedding (t-SNE)

- An unsupervised, randomized algorithm, **used only for visualization**
- Applies a non-linear dimensionality reduction technique where the focus is on **keeping the very similar data points close together in lower-dimensional space**.
- Preserves the local structure of the data using student t-distribution to compute the similarity between two points in lower-dimensional space.



Perplexity=5



perplexity=30



Perplexity=40



```
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.manifold import TSNE

# Load dataset
iris = datasets.load_iris()
X = iris.data
y = iris.target

# Initialize t-SNE
tsne = TSNE(n_components=2, random_state=0)

# Run t-SNE and get the transformed 2D representation
X_2d = tsne.fit_transform(X)
```


```
# Create a scatter plot
plt.figure(figsize=(8, 6))

# Scatter plot for each class label
for i, label in enumerate(iris.target_names):
    plt.scatter(X_2d[y == i, 0], X_2d[y == i, 1], label=label)

plt.legend()
plt.xlabel('t-SNE feature 1')
plt.ylabel('t-SNE feature 2')
plt.title('2D t-SNE on Iris Dataset')

# Show the plot
plt.show()
```

2D t-SNE on Iris Dataset



Uniform Manifold Approximation and Projection (UMAP)

UMAP has several advantages:


- **Speed:** UMAP is generally faster than other techniques like t-SNE, which makes it scalable to larger datasets.
 - **Preservation of Global and Local Structure:** While t-SNE focuses on preserving local structure, UMAP maintains both the local and global aspects of data.
-

UMAP has several advantages:

- **Compatibility**: UMAP can be used in a wide range of applications, not just for visualization. It can be used for clustering, anomaly detection, and more.
 - **Less Arbitrary Parameters**: UMAP has fewer parameters to tune, and they are often easier to interpret (like "minimum distance" and "number of neighbors").
-

UMAP has several advantages:

- **Theoretical Framework:** UMAP comes with a solid mathematical foundation, based on Riemannian geometry and algebraic topology.



```
import matplotlib.pyplot as plt
from sklearn import datasets
import umap

# Load dataset
iris = datasets.load_iris()
X = iris.data
y = iris.target

# Initialize UMAP
umap_model = umap.UMAP(n_neighbors=15, min_dist=0.1,
n_components=2)

# Run UMAP and get the transformed 2D representation
X_2d = umap_model.fit_transform(X)
```


```
# Create a scatter plot
plt.figure(figsize=(8, 6))

# Scatter plot for each class label
for i, label in enumerate(iris.target_names):
    plt.scatter(X_2d[y == i, 0], X_2d[y == i, 1], label=label)


plt.legend()
plt.xlabel('UMAP feature 1')
plt.ylabel('UMAP feature 2')
plt.title('2D UMAP on Iris Dataset')

# Show the plot
plt.show()
```

2D UMAP on Iris Dataset



2D t-SNE on Iris Dataset



2D UMAP on Iris Dataset

