




SVMs, Decision Trees, Random Forest, and *k*-NNs

Machine Learning for Engineering Applications


Fall 2023

Support Vector Machines (SVM)


- **SVM** is one of the most power and most used classifier in the world!
 - The Python community has created many SVM libraries that support all types of classification problems
 - Non-linear classification is the essential part of this tool
-



Which hyperplane?



- **SVM** determines many parts of a classification
- Identifying the vectors
- Drawing the hyperplanes
- Calculate direction (magnitude)
- Margins
- The distance of the hyperplanes
- So on...
- **Reason:** The best optimized (*larger*) boundary will help the model not to overfit



- **SVM** determines many parts of a classification

$$w_0 + \mathbf{w}^T \mathbf{x}_{pos} = 1$$

$$\Rightarrow \mathbf{w}^T (\mathbf{x}_{pos} - \mathbf{x}_{neg}) = 2$$

$$w_0 + \mathbf{w}^T \mathbf{x}_{neg} = -1$$

*Margin
distance*

- Normalize the weights:

$$\|\mathbf{w}\| = \sqrt{\sum_{j=1}^m w_j^2}$$

- You end up with:

$$\frac{w^T (x_{pos} - x_{neg})}{\|w\|} = \frac{2}{\|w\|}$$

Diagram illustrating the margin maximization process:

- The expression $\frac{2}{\|w\|}$ is shown in a box.
- An arrow points from this box to a yellow box labeled "Maximize this part!".
- Another arrow points from the box containing $\frac{2}{\|w\|}$ down to a yellow box labeled "Margin".

- For evaluation, the margin is expressed in its reciprocal:

$$\frac{1}{2} \|w\|^2$$

- **We need a slack variable (ξ):**
- ξ : This drives the model to be a soft-margin classifier
- Hyperplane effect:

$$w_0 + \mathbf{w}^T \mathbf{x}^{(i)} \geq 1 - \xi^{(i)} \quad \text{if } y^{(i)} = 1$$


$$w_0 + \mathbf{w}^T \mathbf{x}^{(i)} \leq -1 + \xi^{(i)} \quad \text{if } y^{(i)} = -1$$

$$\text{for } i = 1 \dots N$$


- You now want to minimize the ξ :

$$\frac{1}{2} \|\mathbf{w}\|^2 + C \left(\sum_i \xi^{(i)} \right)$$

- The C-variable controls the penalty intensity
 - Large C-values: Large Error Penalties / Vice-Versa
-




Large value for
parameter C



Small value for
parameter C

SVM LEAP

Demo



SVM with Non- Linearity

- The XOR non-linear example with random noise


```
import matplotlib.pyplot as plt
import numpy as np
```

```
np.random.seed(1)
```

```
X_xor = np.random.randn(200, 2)
```


```
y_xor = np.logical_xor(X_xor[:, 0] > 0,  
                        ... X_xor[:, 1] > 0)
```

```
y_xor = np.where(y_xor, 1, -1)
```





- Kernel methods, or Kernel SVM, helps with this type of non-linear issue
- The kernel method is a transformation of the data: 2D to a 3D space:

$$\phi(x_1, x_2) = (z_1, z_2, z_3) = (x_1, x_2, x_1^2 + x_2^2)$$




$\phi \rightarrow$






$$z_3 \xrightarrow{\phi^{-1}} x_2$$



- The transformation:

$$\mathbf{x}^{(i)T} \mathbf{x}^{(j)}$$



$$\mathcal{K}(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \phi(\mathbf{x}^{(i)})^T \phi(\mathbf{x}^{(j)})$$

$$\mathcal{K}(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \exp\left(-\frac{\|\mathbf{x}^{(i)} - \mathbf{x}^{(j)}\|^2}{2\sigma^2}\right)$$




Radial Basis Function (RBF)


$$\mathcal{K}(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \exp\left(-\gamma \|\mathbf{x}^{(i)} - \mathbf{x}^{(j)}\|^2\right)$$


$$\gamma = \frac{1}{2\sigma^2}$$

Low Gamma Values



Large Gamma Values




- The XOR non-linear example with random noise

```
svm = SVC(kernel='rbf', random_state=1, gamma=0.10, C=10.0)
```

```
svm.fit(X_xor, y_xor)
```

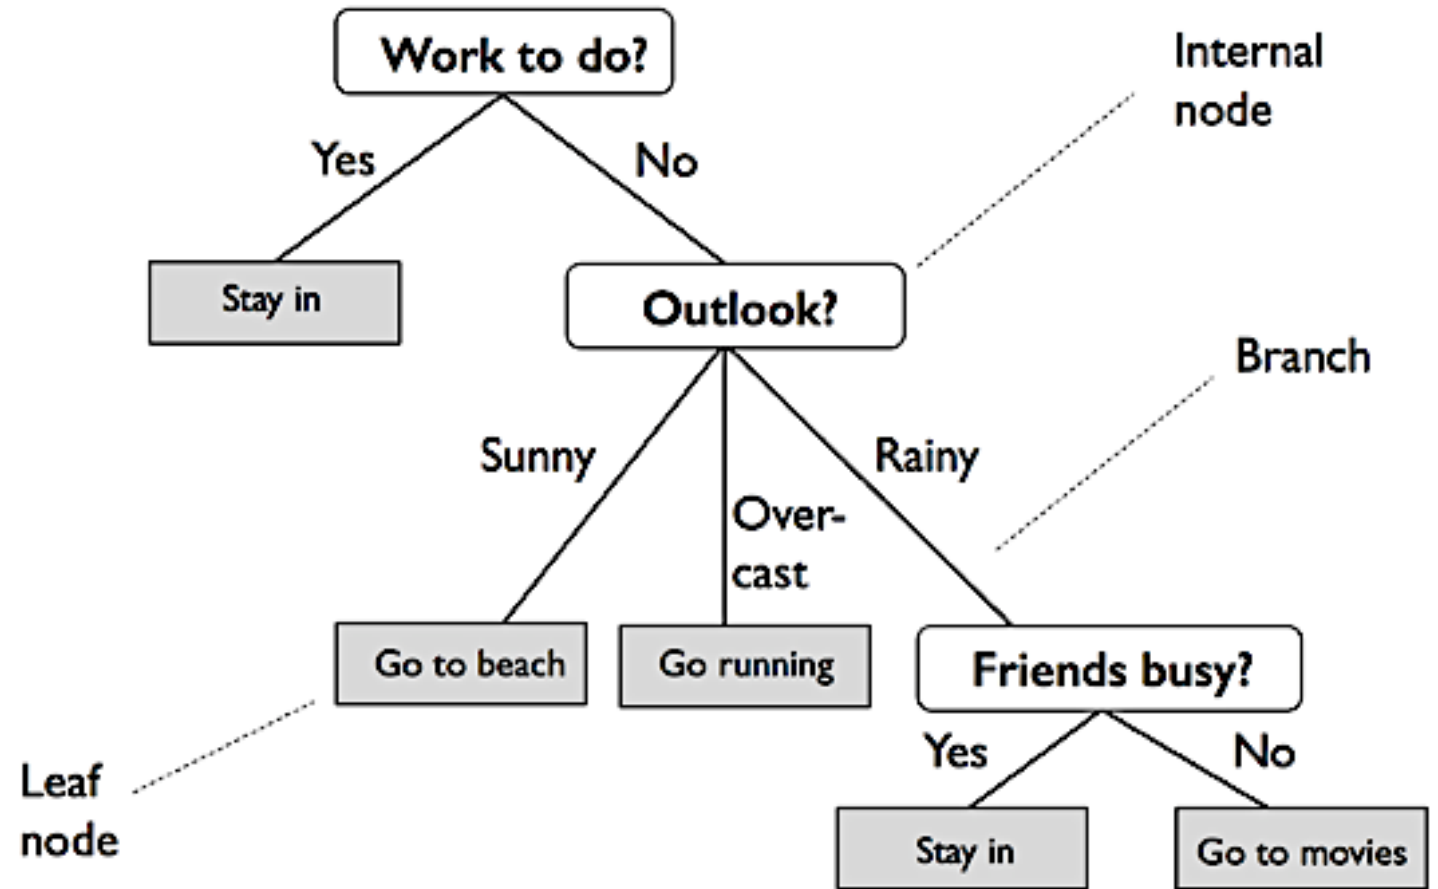
```
plot_decision_regions(X_xor, y_xor, classifier=svm,  
                      ...test_idx=range(y_train.size,  
                                         ..y_train.size + y_test.size))
```

```
plt.legend(loc='upper left')  
plt.show()
```




Decision Tree Learning

- Decision Tree classification
- A classifier that breaks down a decision of which class a new sample belongs to through Q&As



- Decision Tree is a model developed by deducting the class labels of the samples
- The splits are based on features with the largest **Information Gains** (IG)
- Samples @ each node belong to the same class
- Depending on the data: Trees can have a lot of nodes -> **overfitting** -> pruning the tree helps this issue by setting maximum depths (*hyperparameter*)




- What makes the tree split?
- **Information Gains** – To know the gain, there must be a calculation relationship between the **parent node** and the **child nodes**.

- The overall idea:

$$IG(D_p, f) = I(D_p) - \sum_{j=1}^m \frac{N_j}{N_p} I(D_j)$$

- Different between the impurities of the parent node and the aggregated child nodes



- f – feature (column) that performs the split
- I – Impurity function
- D_p – dataset of the parent
- D_j – dataset of the j^{th} child node
- N_p - # of samples @ parent node
- N_j - # of samples @ the j^{th} child node

$$IG(D_p, f) = \underbrace{I(D_p)}_{\text{Parent's}} - \underbrace{\sum_{j=1}^m \frac{N_j}{N_p} I(D_j)}_{\text{Children's}}$$

- **Binary Decision Tree**
- Easier to implement by most ML libraries
- Each Parent node can only spawn 2 child nodes:

$$IG(D_p, f) = I(D_p) - \frac{N_{left}}{N_p} I(D_{left}) - \frac{N_{right}}{N_p} I(D_{right})$$

Diagram illustrating the components of the Information Gain formula:

- Parent's (under $I(D_p)$)
- Child 1 (under $\frac{N_{left}}{N_p} I(D_{left})$)
- Child 2 (under $\frac{N_{right}}{N_p} I(D_{right})$)

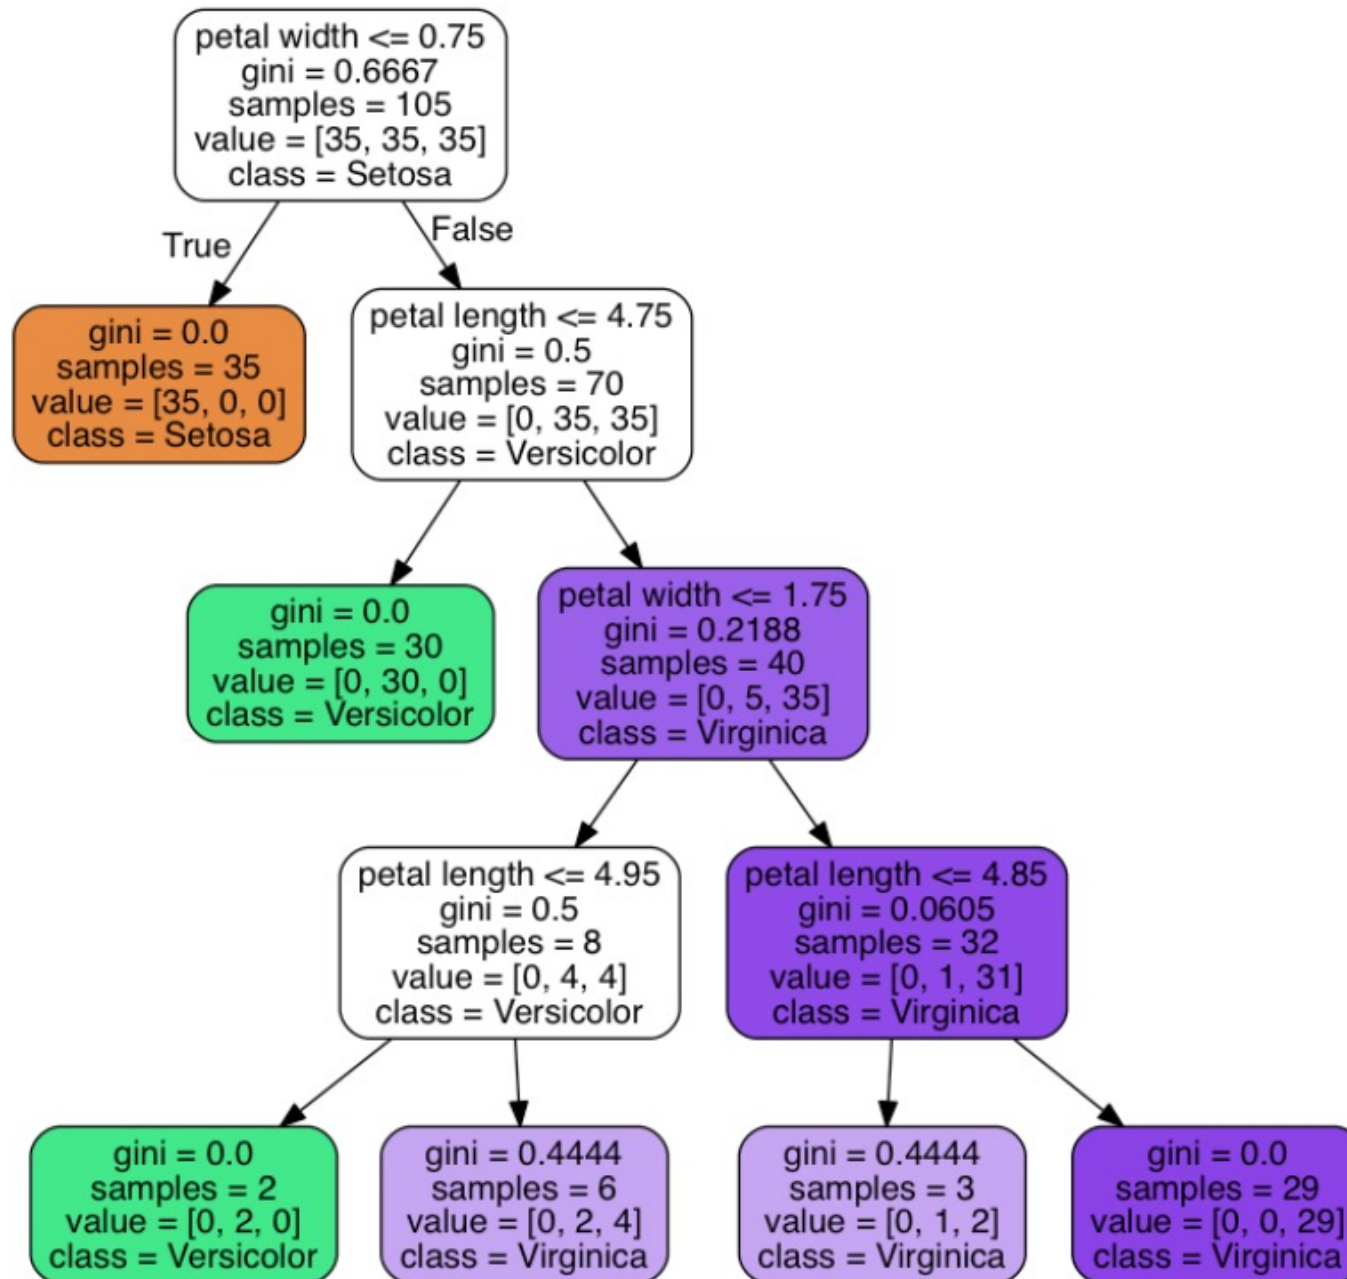
- **Types of Binary Decision Trees:**
 1. Gini
 2. Entropy
 3. Classification Error

```
from sklearn.tree import DecisionTreeClassifier
tree = DecisionTreeClassifier(criterion='gini', max_depth=4,
                              ... random_state=1)

tree.fit(X_train, y_train)
X_combined = np.vstack((X_train, X_test))
y_combined = np.hstack((y_train, y_test))

plot_decision_regions(X_combined,
                      ... y_combined,
                      ... classifier=tree,
                      ... test_idx=range(y_train.size,
                                         .. y_train.size + y_test.size))

plt.xlabel('petal length [cm]')
plt.ylabel('petal width [cm]')
plt.legend(loc='upper left')
plt.show()
```



See textbook
code example
to visualize a
decision tree
(Chapter 3)

Random Forests

- **Random Forest** is random by nature
 - “*Easy*” to implement, hard to explain.
 - Basically: *A collection of Decision Trees*
 - Each Decision Tree is randomly trained
 - **The goal:** Aggregate the accuracy of all the trees
-

- **Implementation of the Random Forest:**

1. Draw a random **bootstrap** sample of size n (randomly choose n samples from the training set with replacement).
 2. Grow a decision tree from the bootstrap sample. At each node:
 - a. Randomly select d features without replacement.
 - b. Split the node using the feature that provides the best split according to the objective function, for instance, maximizing the information gain.
 3. Repeat the steps 1-2 k times.
 4. Aggregate the prediction by each tree to assign the class label by **majority vote**.
-


- **Random Forest Classification:**

```
from sklearn.ensemble import RandomForestClassifier


forest = RandomForestClassifier(criterion='gini',
    ... n_estimators=25,    #25 random decision trees!
    ... random_state=1)
forest.fit(X_train, y_train)
plot_decision_regions(X_combined, y_combined,
    ... classifier=forest, test_idx= range
        ..(y_train.size, y_train.size +
        .. y_test.size))
plt.xlabel('petal length')
plt.ylabel('petal width')
plt.legend(loc='upper left')
plt.show()
```

***k*-Nearest Neighbors**

- k -Nearest Neighbor (kNN) is a supervised **instance-based** learning algorithm
- k – is a variable
- **How does it work:**
 1. Choose the # of k and distance metric
 2. Find the k -NN of the sample that needs to be classified
 3. Assign the category label based on majority vote




- **The Pro:**
 - New data is easily determined by the votes (adaptable)
- **The Con:**
 - The large number of new samples will take over your physical memory (a linear growth)
 - This will depend on the # of features in the dataset



- **k-NN Rules:**

- If the votes are tied:
 - Distance to the closest
- If two or more classes have the **same distance**:
 - Which ever class is listed first in the dataset (*user-determined*)
- The **k-factor** determines great balance, overfitting or underfitting.
- Distance is controlled by the **p-parameter** (*in the code*)



- Popular Distances:

Metrics intended for real-valued vector spaces:

identifier	class name	args	distance function
"euclidean"	EuclideanDistance		$\sqrt{\text{sum}((x - y)^2)}$
"manhattan"	ManhattanDistance		$\text{sum}(x - y)$
"chebyshev"	ChebyshevDistance		$\max(x - y)$
"minkowski"	MinkowskiDistance	p	$\text{sum}(x - y ^p)^{1/p}$
"wminkowski"	WMinkowskiDistance	p, w	$\text{sum}(w * (x - y) ^p)^{1/p}$
"seuclidean"	SEuclideanDistance	V	$\sqrt{\text{sum}((x - y)^2 / V)}$
"mahalanobis"	MahalanobisDistance	V or VI	$\sqrt{(x - y)' V^{-1} (x - y)}$

- p=1 → Manhattan
- p=2 → Euclidean

[illegible]

For next class...

- **Start reading Chapter 4**
 - **HW# 2 will be posted after class**
-