



Logistical Regression and Scikit-Learn

Machine Learning for Engineering Applications

Fall 2023

Scikit-learn



- Website:
 - <https://scikit-learn.org/stable/>
 - Simple and efficient tools for data mining and data analysis
 - Accessible to everybody, and reusable in various contexts
 - Built on **NumPy**, **SciPy**, and **matplotlib**
 - Open source, commercially usable




Classification

Identifying which category an object belongs to.

Applications: Spam detection, image recognition.

Algorithms: SVM, nearest neighbors, random forest, and more...




Examples

Regression

Predicting a continuous-valued attribute associated with an object.

Applications: Drug response, Stock prices.

Algorithms: SVR, nearest neighbors, random forest, and more...




Examples

Clustering

Automatic grouping of similar objects into sets.

Applications: Customer segmentation, Grouping experiment outcomes

Algorithms: k-Means, spectral clustering, mean-shift, and more...




Examples

Dimensionality reduction

Reducing the number of random variables to consider.

Applications: Visualization, Increased efficiency

Algorithms: PCA, feature selection, non-negative matrix factorization, and more...



Model selection

Comparing, validating and choosing parameters and models.

Applications: Improved accuracy via parameter tuning

Algorithms: grid search, cross validation, metrics, and more...




Preprocessing

Feature extraction and normalization.

Applications: Transforming input data such as text for use with machine learning algorithms.

Algorithms: preprocessing, feature extraction, and more...



Logistic Regression (classification)


- Its logistic because of **probabilities**...not binary logic
- This method is used when linear classification does not perform well in creating categories
- Mode of thinking:

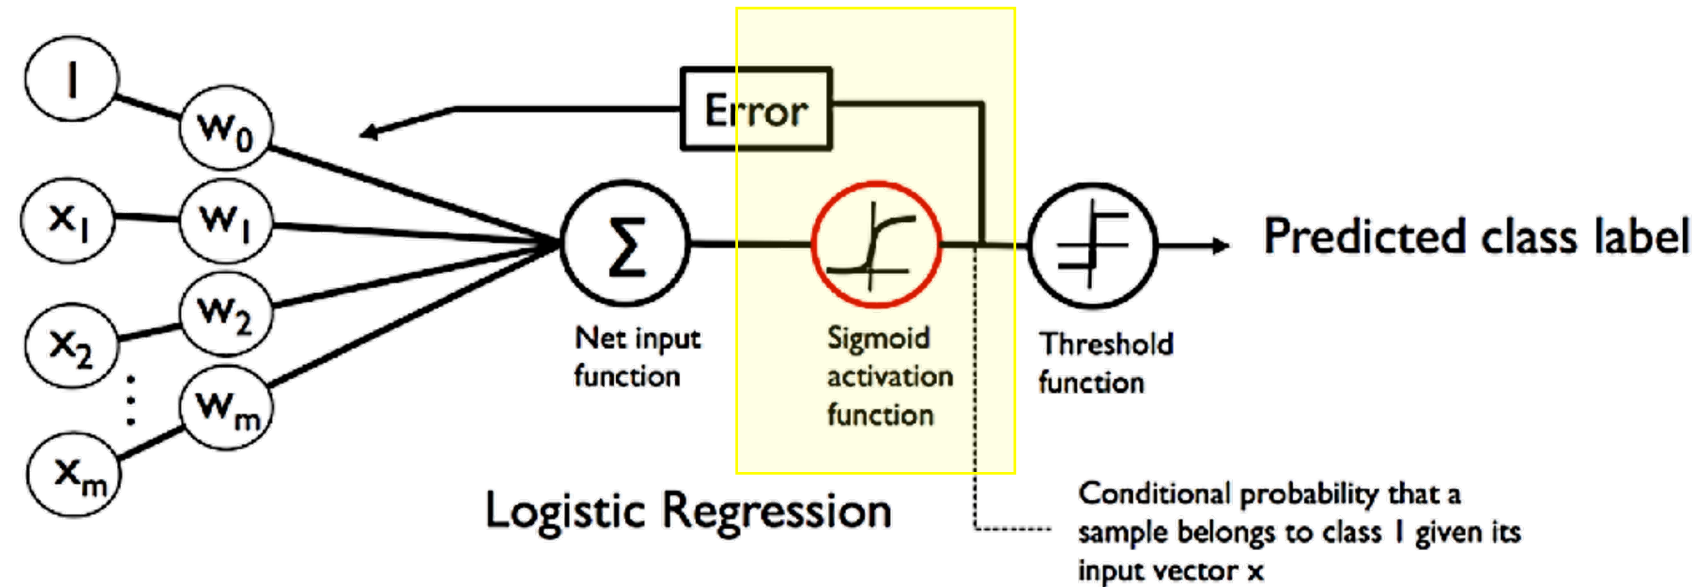
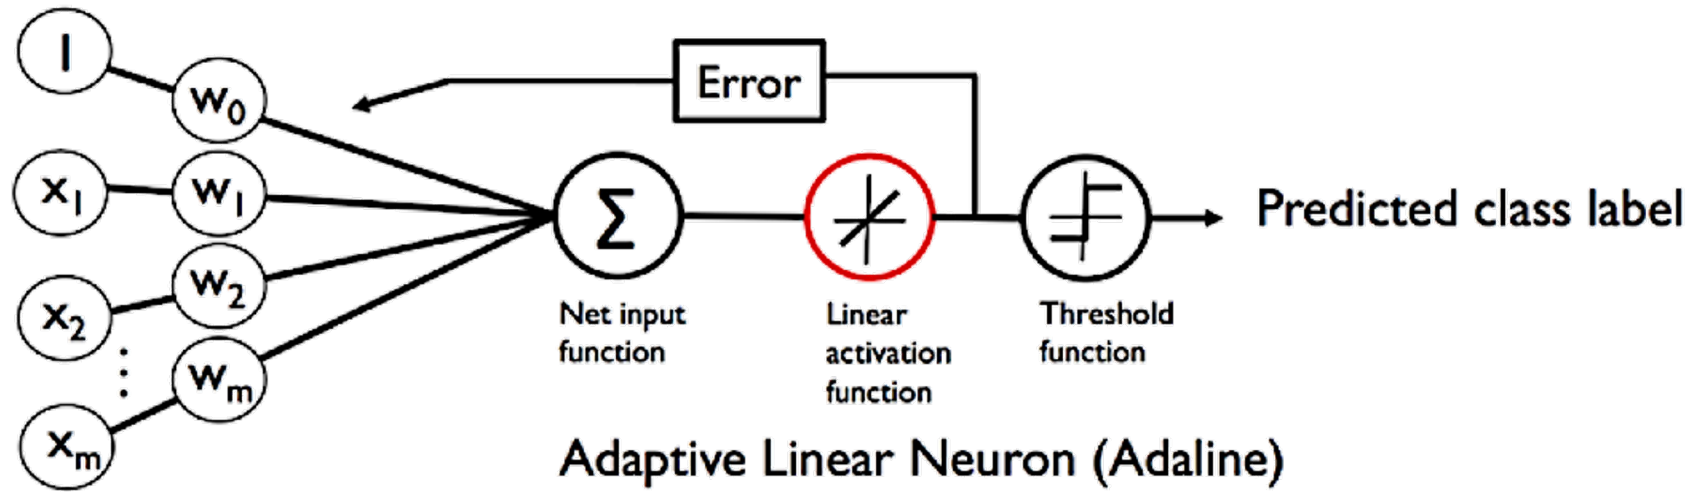
$$\text{logit}(p) = \log \frac{p}{(1-p)}$$

$$\text{logit}(p(y=1 | \mathbf{x})) = w_0x_0 + w_1x_1 + \cdots + w_mx_m = \sum_{i=0}^m w_ix_i = \mathbf{w}^T \mathbf{x}$$

- Goal: Predict the probability of a sample goes into a class
- So, we need the inverse of *logit()*-'
- Inverse: **sigmoid-function**
- **Activation function** is the sigmoid-function



$$\phi(z) = \frac{1}{1 + e^{-z}}$$





- To simplify the decision:

$$\hat{y} = \begin{cases} 1 & \text{if } z \geq 0.0 \\ 0 & \text{otherwise} \end{cases}$$



$$\hat{y} = \begin{cases} 1 & \text{if } \phi(z) \geq 0.5 \\ 0 & \text{otherwise} \end{cases}$$

Logistic Cost Function

- Review, **SSE**: $J(\mathbf{w}) = \sum_i \frac{1}{2} \left(\phi(z^{(i)}) - y^{(i)} \right)^2$

- **Likelihood** (L) is based on a probability value as a function of how well the weights are adjusted in the network:


$$L(\mathbf{w}) = P(\mathbf{y} | \mathbf{x}; \mathbf{w})$$

- Therefore,

$$P(\mathbf{y} | \mathbf{x}; \mathbf{w}) = \prod_{i=1}^n P\left(y^{(i)} | \mathbf{x}^{(i)}; \mathbf{w}\right) = \prod_{i=1}^n \left(\phi(z^{(i)}) \right)^{y^{(i)}} \left(1 - \phi(z^{(i)}) \right)^{1-y^{(i)}}$$

- Maximize likelihood with natural log:

$$l(\mathbf{w}) = \log L(\mathbf{w}) = \sum_{i=1}^n \left[y^{(i)} \log(\phi(z^{(i)})) + (1 - y^{(i)}) \log(1 - \phi(z^{(i)})) \right]$$



- Therefore, the **cost function**:

$$J(\mathbf{w}) = \sum_{i=1}^n \left[-y^{(i)} \log(\phi(z^{(i)})) - (1 - y^{(i)}) \log(1 - \phi(z^{(i)})) \right]$$

, and it must be deterministic as the sigmoid decision:

$$J(\phi(z), y; \mathbf{w}) = \begin{cases} -\log(\phi(z)) & \text{if } y = 1 \\ -\log(1 - \phi(z)) & \text{if } y = 0 \end{cases}$$

- Weights get updated with gradient descent:


$$w_j := w_j + \eta \sum_{i=1}^n \left(y^{(i)} - \phi(z^{(i)}) \right) x_j^{(i)}$$

Adaline

$$\Delta w_j = -\eta \frac{\partial J}{\partial w_j}$$

Logistic Regression

$$J(\mathbf{w}) = \sum_{i=1}^n \left[-y^{(i)} \log(\phi(z^{(i)})) - (1 - y^{(i)}) \log(1 - \phi(z^{(i)})) \right]$$



Over & Under Fitting

- **Overfitting** a model is a very common issue in mathematics, statistics, and Machine Learning
 - Model does well with the training data
 - Model “*fails*” with the test data
 - ***Reasons***: many
 - Too many parameters, or
 - Model too complex for the parameters, or
 - Quality of data is bad, or
 - Not enough data, or
 - Regulation of weights, or
 - Etc....
-


- **Underfitting:** the model is not sophisticated enough to handle the number of parameters from the dataset
 - Model will have to scarifies precision with very “loose” decisions
 - **Fix:**
 - Change the type of model, or
 - Reduce the # of features, or
 - Quality of data is bad, or
 - Regulation of weights, or
 - Etc....
-

- The **L1 & L2 Regularization** techniques help model to not over (or under) fit
- L2 regularization is defined as:

$$\frac{\lambda}{2} \|\mathbf{w}\|^2 = \frac{\lambda}{2} \sum_{j=1}^m w_j^2$$

- Lambda is the regularization parameter (*Dr. Valles's terms: the magnitude of the penalty*)
 - The L2 helps a cost function to level the playing field of the weights -> *preventing overfitting* -> *model learns in a more flexible manner*.
-


- The effect in the *Logical Regression Cost Function*:

$$J(\mathbf{w}) = \sum_{i=1}^n \left[-y^{(i)} \log(\phi(z^{(i)})) - (1 - y^{(i)}) \log(1 - \phi(z^{(i)})) \right] + \frac{\lambda}{2} \|\mathbf{w}\|^2$$


- **Lambda** -> increase values -> the regularization strength increases
- In Python code: $\lambda = \frac{1}{C}$
- **C-parameter**: Inverse Regularization Parameter

```
weights, params = [], []
for c in np.arange(-5, 5):
    ... lr = LogisticRegression(C=10.**c, random_state=1)
    ... lr.fit(X_train_std, y_train)
    ... weights.append(lr.coef_[1])
    ... params.append(10.**c)
weights = np.array(weights)

plt.plot(params, weights[:, 0], label='petal length')
plt.plot(params, weights[:, 1], linestyle='--',
    ... label='petal width')
plt.ylabel('weight coefficient')
plt.xlabel('C')
plt.legend(loc='upper left')
plt.xscale('log')
plt.show()
```



Python IDEs

- Jupyter Notebooks <https://jupyter.org/>
 - ATOM: <https://atom.io/>
 - PyCharm: <https://www.jetbrains.com/pycharm/>
 - Visual Studio Code (VS Code): <https://code.visualstudio.com/>
 - Spyder: <https://www.spyder-ide.org/> or though...
 - Anaconda: <https://www.anaconda.com/>
-

For next class...

- **Continue reading Chapters 2 & 3**
 - **Homework #1 will be posted on Wed.**
 - **LEAP Cluster training on Wed.**
-

Hands On

- Open your IDE to develop the in-class exercise:
 - Use the IRIS dataset on the Logistic Regression model
 - **Find the best C-values for the highest accuracy**
 - You can work solo, in pair, or table!
 - Provide the following analysis:
 - Training accuracy
 - Test accuracy
 - Number of test misclassifications
 - **Use the code in Page 56** to plot the decision regions
-