




Deep Convolutional Neural Networks – Part 2

Machine Learning for Engineering Applications

Fall 2023

DCNN + TensorFlow

- **Input layer:** 28x28 pixel resolution images (MNIST – digits)
 - **Batch size:** 28 x 28 x 1 (one channel – gray images)
 - **Kernel size:** 5x5
 - **1st conv:** 32 output feature maps
 - **2nd conv:** 64 output feature maps
 - max-pooling after each conv-layer
 - **Two dense-layers** (1st-1024, 2nd 10 [softmax])
-



- **Tensor dimensions:**

- Input: [batchsize× 28× 28×1]
 - Conv_1: [batchsize× 24× 24×32]
 - Pooling_1: [batchsize×12×12×32]
 - Conv_2: [batchsize×8×8× 64]
 - Pooling_2: [batchsize× 4× 4× 64]
 - FC_1: [batchsize×1024]
 - FC_2 and softmax layer: [batchsize×10]
-

- Loading the MNIST dataset:

```
X_data, y_data = load_mnist('./mnist/', kind='train')
print('Rows: {}, Columns: {}'.format(X_data.shape[0], X_data.shape[1]))
    Rows: 60000, Columns: 784
X_test, y_test = load_mnist('./mnist/', kind='t10k')
print('Rows: {}, Columns: {}'.format(X_test.shape[0], X_test.shape[1]))
    Rows: 10000, Columns: 784
X_train, y_train = X_data[:50000,:], y_data[:50000]
X_valid, y_valid = X_data[50000:,:], y_data[50000:]

print('Training: ', X_train.shape, y_train.shape)
    Training: (50000, 784) (50000,)
print('Validation: ', X_valid.shape, y_valid.shape)
    Validation: (10000, 784) (10000,)
print('Test Set: ', X_test.shape, y_test.shape)
    Test Set: (10000, 784) (10000,)
```

- Data preprocessing

```
from keras.datasets import mnist
from keras.utils import to_categorical

(train_images, train_labels), (test_images, test_labels) =
    ...mnist.load_data()

# training tensor and normalization
train_images = train_images.reshape((60000, 28, 28, 1))
train_images = train_images.astype('float32') / 255

# test tensor and normalization
test_images = test_images.reshape((10000, 28, 28, 1))
test_images = test_images.astype('float32') / 255

# string to numerical
train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)
```

- Straight forward CNN Layers (*think of the MNIST dataset*)

```
from keras import layers
from keras import models
```

```
model = models.Sequential()
```

```
model.add(layers.Conv2D(32, (3, 3), activation='relu',  
                        ...input_shape=(28, 28, 1)))
```

```
model.add(layers.MaxPooling2D((2, 2)))
```

```
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```


```
model.add(layers.MaxPooling2D((2, 2)))
```

```
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

- Print out the summary of your CNN model

```
model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
maxpooling2d_1 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_2 (Conv2D)	(None, 11, 11, 64)	18496
maxpooling2d_2 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_3 (Conv2D)	(None, 3, 3, 64)	36928
Total params: 55,744		
Trainable params: 55,744		
Non-trainable params: 0		

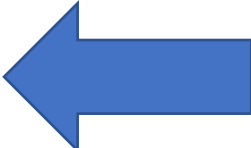


- Adding the dense layers

```
model.add(layers.Conv2D(32, (3, 3), activation='relu',  
                        ...input_shape=(28, 28, 1)))  
model.add(layers.MaxPooling2D((2, 2)))  
  
model.add(layers.Conv2D(64, (3, 3), activation='relu'))  
model.add(layers.MaxPooling2D((2, 2)))  
  
model.add(layers.Conv2D(64, (3, 3), activation='relu'))  
  
model.add(layers.Flatten())  
model.add(layers.Dense(64, activation='relu'))  
model.add(layers.Dense(10, activation='softmax'))
```

- Adding the dense layers

Layer (type)	Output Shape	Param #
=====	=====	=====
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
maxpooling2d_1 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_2 (Conv2D)	(None, 11, 11, 64)	18496
maxpooling2d_2 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_3 (Conv2D)	(None, 3, 3, 64)	36928
flatten_1 (Flatten)	(None, 576)	0
dense_1 (Dense)	(None, 64)	36928
dense_2 (Dense)	(None, 10)	650
=====	=====	=====
Total params: 93,322		
Trainable params: 93,322		
Non-trainable params: 0		



- Optimize / Train / Test

```
model.compile(optimizer='rmsprop',  
              ...loss='categorical_crossentropy', metrics=['accuracy'])  
  
model.fit(train_images, train_labels, epochs=5, batch_size=64)  
  
test_loss, test_acc = model.evaluate(test_images, test_labels)
```

Epoch 1/5

938/938 [=====] - 54s 57ms/step -
loss: 0.1606 - accuracy: 0.9492

Epoch 2/5

938/938 [=====] - 55s 59ms/step -
loss: 0.0457 - accuracy: 0.9857

Epoch 3/5

938/938 [=====] - 53s 57ms/step -
loss: 0.0319 - accuracy: 0.9902

Epoch 4/5

938/938 [=====] - 52s 55ms/step -
loss: 0.0227 - accuracy: 0.9930

Epoch 5/5

938/938 [=====] - 46s 49ms/step -
loss: 0.0172 - accuracy: 0.9945

313/313 [=====] - 2s 8ms/step - loss:
0.0330 - accuracy: 0.9911

Input

+

Convolutional operations

=


Output

Input Image
(28x28 pixels)



Convolutional Layer 1

Filter-Weights
(5x5 pixels)



(14x14 pixels)



(16 channels)

Convolutional Layer 2

Filter-Weights
(5x5 pixels)




(7x7 pixels)




(36 channels)

Fully-Connected
Layer



(128 features)

Output
Layer



(10 features)

Class

0
1
2
3
4
5
6
7
8
9

7

CATS vs.
DOGS

- Dataset: <https://www.kaggle.com/c/dogs-vs-cats/data>



- How to deal with images from other sources?
 - How do you organize and format the images to your needs?
 - How do you split training / validation / testing?
 - Sometimes:
 - The longer and organized the code is, the less likely to make a mistake....just comment a lot.
-

```
import os, shutil
```

```
original_dataset_dir = '/home/dvalles/Downloads/kaggle_original_data'
```

```
base_dir = '/home/dvalles/Downloads/cats_and_dogs_small'
```

```
os.mkdir(base_dir)
```

```
train_dir = os.path.join(base_dir, 'train')  
os.mkdir(train_dir)
```

```
validation_dir = os.path.join(base_dir, 'validation')  
os.mkdir(validation_dir)
```

```
test_dir = os.path.join(base_dir, 'test')  
os.mkdir(test_dir)
```

Directories for
the training /
validation / test
splits

```
train_cats_dir = os.path.join(train_dir, 'cats')  
os.mkdir(train_cats_dir)
```

```
train_dogs_dir = os.path.join(train_dir, 'dogs')  
os.mkdir(train_dogs_dir)
```

Own training
folders for cats
and dogs

```
validation_cats_dir = os.path.join(validation_dir, 'cats')  
os.mkdir(validation_cats_dir)
```

```
validation_dogs_dir = os.path.join(validation_dir, 'dogs')  
os.mkdir(validation_dogs_dir)
```

Own
validation
folders for
cats and
dogs

```
test_cats_dir = os.path.join(test_dir, 'cats')  
os.mkdir(test_cats_dir)
```

```
test_dogs_dir = os.path.join(test_dir, 'dogs')  
os.mkdir(test_dogs_dir)
```

Own test
folders for
cats and
dogs

```
fnames = ['cat.{}.jpg'.format(i) for i in range(1000)]  
for fname in fnames:  
    src = os.path.join(original_dataset_dir, fname)  
    dst = os.path.join(train_cats_dir, fname)  
    shutil.copyfile(src, dst)
```

1st 1000 cat
copies to the
training folder for
cat

```
fnames = ['cat.{}.jpg'.format(i) for i in range(1000, 1500)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
    dst = os.path.join(validation_cats_dir, fname)
    shutil.copyfile(src, dst)
```

The next 500 cat
copies to the
validation folder
for cat

```
fnames = ['cat.{}.jpg'.format(i) for i in range(1500, 2000)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
    dst = os.path.join(test_cats_dir, fname)
    shutil.copyfile(src, dst)
```

The next 500 cat
copies to the test
folder for cat

```
fnames = ['dog.{}.jpg'.format(i) for i in range(1000)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
    dst = os.path.join(train_dogs_dir, fname)
    shutil.copyfile(src, dst)
```

1st 1000 dog copies to
the training folder for
cat

```
fnames = ['dog.{}.jpg'.format(i) for i in range(1000, 1500)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
    dst = os.path.join(validation_dogs_dir, fname)
    shutil.copyfile(src, dst)
```

The next 500 dog
copies to the
validation folder
for cat

```
fnames = ['dog.{}.jpg'.format(i) for i in range(1500, 2000)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
    dst = os.path.join(test_dogs_dir, fname)
    shutil.copyfile(src, dst)
```

The next 500 dog
copies to the test
folder for cat

```
>>> print('total training cat images:', len(os.listdir(train_cats_dir)))  
total training cat images: 1000
```

```
>>> print('total training dog images:', len(os.listdir(train_dogs_dir)))  
total training dog images: 1000
```

```
>>> print('total validation cat images:', len(os.listdir(validation_cats_dir)))  
total validation cat images: 500
```

```
>>> print('total validation dog images:', len(os.listdir(validation_dogs_dir)))  
total validation dog images: 500
```

```
>>> print('total test cat images:', len(os.listdir(test_cats_dir)))  
total test cat images: 500
```

```
>>> print('total test dog images:', len(os.listdir(test_dogs_dir)))  
total test dog images: 500
```

Preprocess the images

- Read the picture files.
 - Decode the JPEG content to RGB grids of pixels.
 - Convert these into floating-point tensors.
 - Rescale the pixel values (between 0 and 255) to the $[0, 1]$ interval
-

- Preprocess the images

```
from tf.keras.preprocessing.image import ImageDataGenerator
```

```
train_datagen = ImageDataGenerator(rescale=1./255)
```

```
test_datagen = ImageDataGenerator(rescale=1./255)
```

```
train_generator = train_datagen.flow_from_directory(  
    train_dir,  
    target_size=(150, 150)  
    batch_size=32,  
    class_mode='binary' )
```

```
validation_generator = test_datagen.flow_from_directory(  
    validation_dir,  
    target_size=(150, 150) ,  
    batch_size=32,  
    class_mode='binary' )
```

- The CNN:

```
from tf.keras import layers
from tf.keras import models
```

```
model = models.Sequential()
```

```
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
```

```
model.add(layers.Flatten())
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

- Optimize / Train & Save

```
from tf.keras import optimizers
```

```
model.compile(loss='binary_crossentropy',  
              ...optimizer=optimizers.RMSprop(lr=1e-4), metrics=['acc'])
```

```
history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=100, #number of gradient step before next epoch  
    epochs=30,  
    validation_data=validation_generator,  
    validation_steps=50)  
    # how many batches to draw from the validation generator for evaluation
```

```
model.save('cats_and_dogs_small_1.h5')
```


- Plot the accuracy and losses

```
import matplotlib.pyplot as plt


acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)



plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```

Training and validation accuracy



Training and validation loss


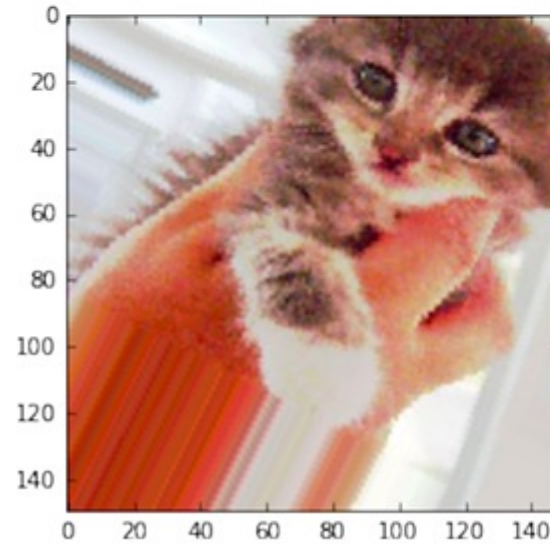






- How to address the problem?
- **Data Augmentation** is a way to compensate the dataset
- Dropout layer were not implemented

- Data Augmentation

```
datagen = ImageDataGenerator(  
    rotation_range=40,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,  
    fill_mode='nearest')
```



- Data Augmentation

rotation_range is a value in degrees (0–180), a range within which to randomly rotate pictures.

width_shift and *height_shift* are ranges (as a fraction of total width or height) within which to randomly translate pictures vertically or horizontally.

shear_range is for randomly applying shearing transformations.

- Data Augmentation

zoom_range is for randomly zooming inside pictures.

horizontal_flip is for randomly flipping half the images horizontally—relevant when there are no assumptions of horizontal asymmetry (for example, real-world pictures).

fill_mode is the strategy used for filling in newly created pixels, which can appear after a rotation or a width/height shift.

• Data Preprocessing

```
train_datagen = ImageDataGenerator( rescale=1./255,rotation_range=40,  
    ...width_shift_range=0.2, height_shift_range=0.2,  
    ...shear_range=0.2, zoom_range=0.2, horizontal_flip=True)
```

```
test_datagen = ImageDataGenerator(rescale=1./255)
```

```
train_generator = train_datagen.flow_from_directory(  
    train_dir, target_size=(150, 150),  
    batch_size=32, class_mode='binary')
```

```
validation_generator = test_datagen.flow_from_directory(  
    validation_dir, target_size=(150, 150),  
    batch_size=32, class_mode='binary')
```

• Dropout Layers

```
model = models.Sequential()
```

```
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)))
```

```
model.add(layers.MaxPooling2D((2, 2)))
```

```
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

```
model.add(layers.MaxPooling2D((2, 2)))
```

```
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
```

```
model.add(layers.MaxPooling2D((2, 2)))
```

```
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
```

```
model.add(layers.MaxPooling2D((2, 2)))
```

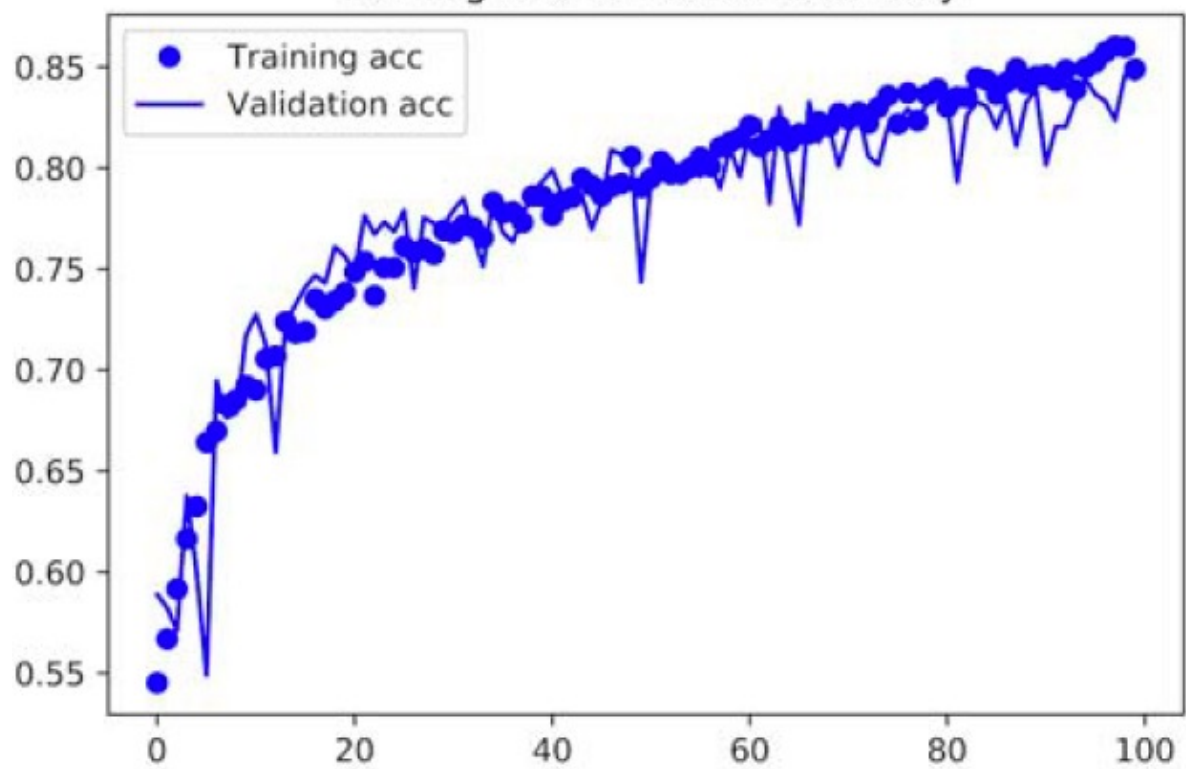
```
model.add(layers.Flatten())
```

```
model.add(layers.Dropout(0.5)) #Try one position, add more if need it
```

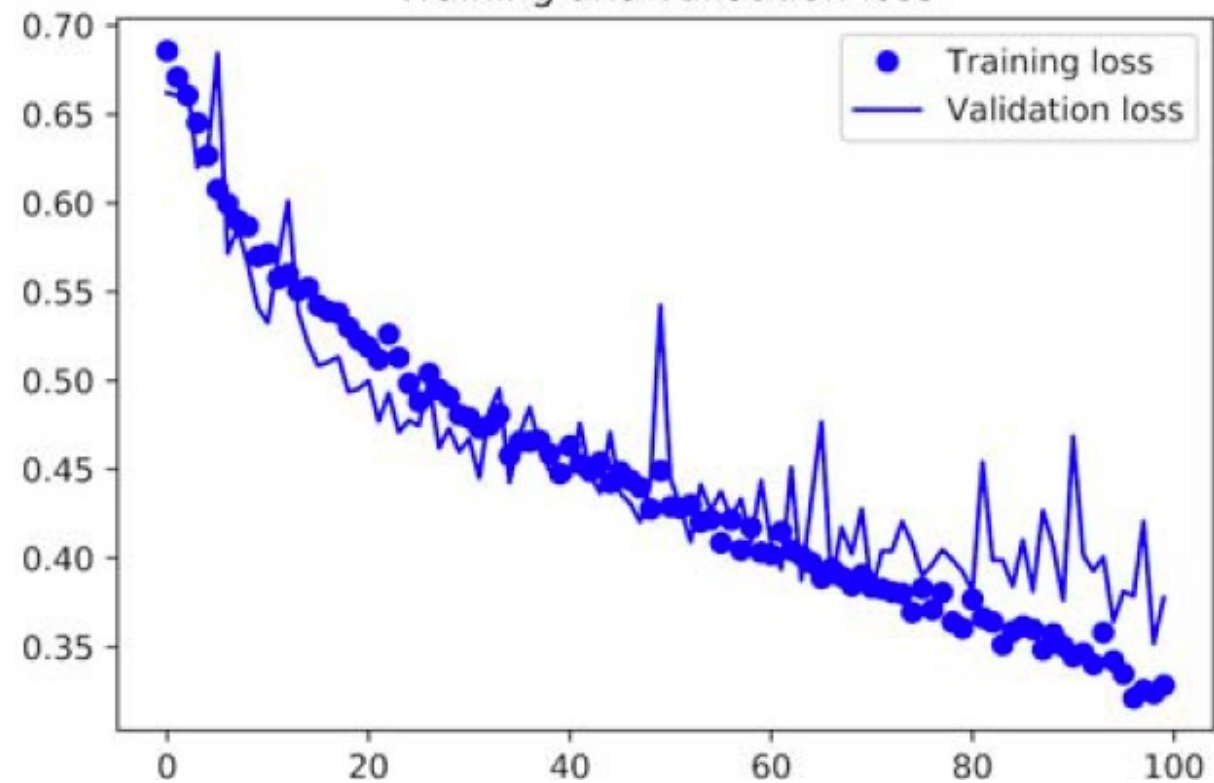
```
model.add(layers.Dense(512, activation='relu'))
```

```
model.add(layers.Dense(1, activation='sigmoid'))
```

Training and validation accuracy



Training and validation loss



MNIST CNN
